

# Feed forward network implementation

1. Implementing Feed forward neural networks with Keras and TensorFlow a. Import the necessary packages b. Load the training and testing data (MNIST/CIFAR10) c. Define the network architecture using Keras d. Train the model using SGD e. Evaluate the network f. Plot the training loss and accuracy for this practical and : code: #a. Import the necessary packages import tensorflow as tf from tensorflow import keras from tensorflow.keras.preprocessing.image import ImageDataGenerator import matplotlib.pyplot as plt import numpy as np #b. Load the training and testing data (MNIST/CIFAR10 train_data_dir = 'mnist-jpg/train' test_data_dir = 'mnist-jpg/test' # Image data generator for training data train_datagen = ImageDataGenerator( rescale=1.0/255.....	3
2. 1. what is mnist 2.x_train, y_train = train_generator[0] x_test, y_test = test_generator[0] why [0] 3.what is this why is this ? ✓ Flatten → converts 28×28 image → 784-dim vector ✓ Dense(128) → 128 neurons, activation = ReLU ✓ Dense(10) → 10 output classes, softmax gives probabilities This is a Feed-Forward Neural Network (no convolution).	
4.what is sgd output related doubt also give me.....	11
3. What is a feed-forward neural network? Why do we use the Flatten layer? Difference between SGD and Adam optimizer. What is the role of the activation function? Why do we use ReLU in hidden layers? What is softmax activation used for? Why do we normalize images? What is the difference between training loss and validation loss? What is overfitting? How do you detect it from graphs? What is categorical cross-entropy? ✓ Program / Practical Questions What will happen if we remove the Flatten layer? Why is the output layer size = 10 in MNIST? What happens if batch size is very large? Why do we use grayscale mode for MNIST? How does ImageDataGenerator help during training? What is one-hot encoding and why is it needed? What change.....	15
4. in code where is sgd used and how that firmula is used.....	20
5. what is vanishing graident problem.....	24
6. if practical is about feed forwards neural network theen how can we back propagate.....	25
7. so to summarize this just tell me how should i explain mycode practical to exteral thatwhat i my code aout.....	27
8. and what does sgd do.....	29
9. ok now #a. Import the necessary packages import tensorflow as tf from tensorflow import keras from tensorflow.keras import layers from tensorflow.keras.layers import Input, Conv2D, Dense, Flatten, Dropout from tensorflow.keras.layers import GlobalMaxPooling2D, MaxPooling2D from tensorflow.keras.layers import BatchNormalization from tensorflow.keras.preprocessing.image import ImageDataGenerator from tensorflow.keras.models import Model from keras.utils import to_categorical import matplotlib.pyplot as plt import numpy as np #b. Load the training and testing data (MNIST/CIFAR10) train_data_dir = 'cifar-10-img/train' test_data_dir = 'cifar-10-img/test' # Image data generator for training data train_datagen = ImageDataGenerator( rescale.....	30
10. explain me this code in three parts for viva first : function by function which does what , important keywords and libraries and third possible questions asked maybe theory based or practical based sppu deep learnig context.....	38
11. viva questions with answers.....	45
12. What is a CNN? Why do we use Conv2D layers? What is the purpose of BatchNormalization? Why use MaxPooling? What is the difference between RGB and grayscale images? What is ReLU and why is it used? Why do we use softmax in the output layer? What is Dropout? Why do we use it? What is sparse categorical cross-entropy? What is overfitting and how do graphs show it? ✓ Practical Viva Questions Why did you resize to 32×32? Why use class_mode='sparse'? Why batch size = 20,000? Why use three Conv blocks? What happens if we remove Dropout? Why do we normalize images (rescale=1/255)? What does model.compile() do? What is the role of SGD? What is the effect of increasing epochs? Why do we use Flatten before Dense? ✓ Advanced Viva Q.....	50
13. #3Build the Image classification model by dividing the model into following 4 stages: a. Loading and preprocessing the image data b. Defining the model's architecture c. Training the model d. Estimating the model's performance import numpy as np import tensorflow as tf from tensorflow.keras.datasets import mnist from tensorflow.keras.utils import to_categorical from tensorflow.keras import layers, models from tensorflow.keras.preprocessing.image import ImageDataGenerator import matplotlib.pyplot as plt #a. Loading and preprocessing the image data train_data_dir = 'mnist-jpg/train' test_data_dir = 'mnist-jpg/test' # Image data generator for training data train_datagen = ImageDataGenerator( rescale=1.0/255 ) # Image data generator f.....	55
14. to summaruze the practical to external explain me in short.....	65

15. ""4ECG Use Autoencoder to implement anomaly detection. Build the model by using: a. Import required libraries b. Upload / access the dataset c. Encoder converts it into latent representation d. Decoder networks convert it back to the original input e. Compile the models with Optimizer, Loss, and Evaluation Metrics #Aa. Import required libraries import numpy as np import pandas as pd import tensorflow as tf from sklearn.model\_selection import train\_test\_split from sklearn.preprocessing import StandardScaler from sklearn.metrics import confusion\_matrix, classification\_report from tensorflow.keras import layers, models import matplotlib.pyplot as plt #B Upload / access the dataset # Load the ECG dataset ecg\_dataset = pd.read\_csv("ecg.csv") ..... .66
16. "" Use Autoencoder to implement anomaly detection. Build the model by using: a. Import required libraries b. Upload / access the dataset c. Encoder converts it into latent representation d. Decoder networks convert it back to the original input e. Compile the models with Optimizer, Loss, and Evaluation Metrics "" #a.Import required libraries import numpy as np import pandas as pd import tensorflow as tf from sklearn.model\_selection import train\_test\_split from sklearn.preprocessing import StandardScaler from sklearn.metrics import confusion\_matrix, classification\_report import matplotlib.pyplot as plt from tensorflow.keras import layers, models #b. Upload / access the dataset dataset = pd.read\_csv("creditcard.csv") # Preproces..... .78
17. 5. Implement the Continuous Bag of Words (CBOW) Model. Stages can be: a. Data preparation b. Generate training data c. Train model d. Output import tensorflow as tf from tensorflow.keras.models import Sequential from tensorflow.keras.layers import Dense,\ Embedding, Lambda from tensorflow.keras.preprocessing.text import Tokenizer import numpy as np import matplotlib.pyplot as plt from sklearn.decomposition import PCA import re data = """We are about to study the idea of a computational process. Computational processes are abstract beings that inhabit computers. As they evolve, processes manipulate other abstract things called data. The evolution of a process is directed by a pattern of rules called a program. People create pr..... .89
18. 6. Object detection using Transfer Learning of CNN architectures a. Load in a pre-trained CNN model trained on a large dataset b. Freeze parameters (weights) in model's lower convolutional layers c. Add custom classifier with several layers of trainable parameters to model d. Train classifier layers on training data available for task e. Fine-tune hyper parameters and unfreeze more layers as needed dataset: CIFAR import tensorflow as tf from tensorflow import keras from tensorflow.keras.applications import VGG16 from tensorflow.keras.models import Model from tensorflow.keras.layers import Dense, Flatten from tensorflow.keras.optimizers import Adam from tensorflow.keras.preprocessing.image import ImageDataGenerator import numpy as ..... .103
19. 6. Object detection using Transfer Learning of CNN architectures a. Load in a pre-trained CNN model trained on a large dataset b. Freeze parameters (weights) in model's lower convolutional layers c. Add custom classifier with several layers of trainable parameters to model d. Train classifier layers on training data available for task e. Fine-tune hyper parameters and unfreeze more layers as needed dataset: Caltech import tensorflow as tf from tensorflow import keras from tensorflow.keras.applications import VGG16 from tensorflow.keras.models import Model from tensorflow.keras.layers import Dense, Flatten from tensorflow.keras.optimizers import Adam from tensorflow.keras.preprocessing.image import ImageDataGenerator import numpy as n..... .114



## 1. Implementing Feed forward neural networks with Keras and TensorFlow

- a. Import the necessary packages
- b. Load the training and testing data (MNIST/CIFAR10)
- c. Define the network architecture using Keras
- d. Train the model using SGD
- e. Evaluate the network
- f. Plot the training loss and accuracy

for this practical and :

code:

```
#a. Import the necessary packages
```

```
import tensorflow as tf
```

```
from tensorflow import keras
```

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
#b. Load the training and testing data (MNIST/CIFAR10
```

```
train_data_dir = 'mnist-jpg/train'
```

```
test_data_dir = 'mnist-jpg/test'
```

```
# Image data generator for training data
```

```
train_datagen = ImageDataGenerator(
```

```
rescale=1.0/255
```

```
)
```

```
# Image data generator for testing data
```

```
test_datagen = ImageDataGenerator(
```

```
rescale=1.0/255
```

```
)
```

```
# Create data generators
```

```
train_batch_size = 10000
```

```
train_generator = train_datagen.flow_from_directory(
```

```
    train_data_dir,
```

```
    target_size=(28, 28), # Resize images to 28x28
```

```
    batch_size=train_batch_size,
```

```
    class_mode='categorical',
```

```
    color_mode='grayscale',# Use 'categorical' for one-hot encoded labels
```

```
    shuffle=True,
```

```
)
```

```
# Load test data without labels (class_mode=None)
```

```
test_batch_size = 2000
```

```
test_generator = test_datagen.flow_from_directory(  
    test_data_dir,  
    target_size=(28, 28), # Resize images to 28x28  
    batch_size=test_batch_size,  
    class_mode='categorical', # Use 'categorical' for one-hot encoded labels  
    color_mode='grayscale',  
    shuffle=True,  
)
```

```
#Selecting first batch containing 10000 images  
x_train, y_train = train_generator[0]  
x_test, y_test = test_generator[0]
```

#c. Define the network architecture using Keras

```
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28, 1)),
    keras.layers.Dense(128, activation="relu"),
    keras.layers.Dense(10, activation="softmax")
])
```

```
#d.Train the model using SGD  
model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model.fit(x_train, y_train, epochs=8, validation_data=(x_test, y_test))
```

#### #e. Evaluate the network

```
test_loss, test_acc = model.evaluate(x_test, y_test)
print("Loss: ", test_loss)
print("Accuracy: ", test_acc)
```

```
n = 20  
plt.imshow(x_test[n])  
predicted_value = model.predict(x_test)  
print("Actual Number: ",np.argmax(y_test[n]))  
print("Predicted Number: ", np.argmax(predicted_value[n]))
```

```
history = history.history  
history.keys()
```

```
plt.plot(history['accuracy'])
plt.plot(history['val_accuracy'])
plt.title("Model Accuracy")
plt.ylabel('accuracy')
```

```
plt.xlabel('epoch')
plt.legend(['Train', "Validation"], loc='upper left')
```

```
plt.plot(history['loss'])
plt.plot(history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
```

```
plt.plot(history['accuracy'])
plt.plot(history['val_accuracy'])
plt.plot(history['loss'])
plt.plot(history['val_loss'])
plt.title('Training Loss and accuracy')
plt.ylabel('accuracy/Loss')
plt.xlabel('epoch')
plt.legend(['accuracy', 'val_accuracy', 'loss', 'val_loss'])
```

explain me this code in three parts for viva first : function by function which does what , important keywords and libraries and third possible questions asked maybe theory based or practical based sppu deep learnig context

Below is a **perfect viva-ready explanation** of your code separated into **3 parts**, exactly as you asked:

## ✓ PART 1 — Function-by-Function Explanation (Line-by-Line What Each Block Does)

### a. Import the necessary packages

python

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt
import numpy as np
```

- ✓ **tensorflow, keras** → for building & training neural networks
- ✓ **ImageDataGenerator** → loads images from folders & preprocesses them
- ✓ **matplotlib.pyplot** → plotting loss & accuracy
- ✓ **numpy** → numerical operations, arrays

## b. Load the training and testing data

You are loading **offline MNIST JPEG images** stored in folders:

```
python

train_data_dir = 'mnist-jpg/train'
test_data_dir = 'mnist-jpg/test'
```

## ImageDataGenerator

```
python

train_datagen = ImageDataGenerator(rescale=1.0/255)
test_datagen = ImageDataGenerator(rescale=1.0/255)
```

- ✓ **rescale=1/255** → converts pixel values 0–255 → 0–1
- ✓ prevents large input values (helps training)

## Train generator

```
python

train_generator = train_datagen.flow_from_directory(
    train_data_dir,
    target_size=(28, 28),
    batch_size=10000,
    class_mode='categorical',
    color_mode='grayscale',
    shuffle=True
)
```

- ✓ Loads all images at once (batch size = 10000)
- ✓ **target\_size=(28, 28)** → resize images
- ✓ **color\_mode='grayscale'** → MNIST is 1 channel
- ✓ **class\_mode='categorical'** → one-hot encoding

Same for test data.

## Select 1st batch (since entire dataset loaded as one batch)

```
python

x_train, y_train = train_generator[0]
x_test, y_test = test_generator[0]
```

- ✓ Extracts all images & labels into arrays.

## c. Define the Neural Network

```
python
```

```
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28, 1)),
    keras.layers.Dense(128, activation="relu"),
    keras.layers.Dense(10, activation="softmax")
])
```

- ✓ **Flatten** → converts 28×28 image → 784-dim vector
- ✓ **Dense(128)** → 128 neurons, activation = ReLU
- ✓ **Dense(10)** → 10 output classes, softmax gives probabilities

**This is a Feed-Forward Neural Network (no convolution).**

## d. Compile & Train

```
python
```

```
model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(x_train, y_train, epochs=8, validation_data=(x_test, y_test))
```

- ✓ Optimizer = **SGD**
- ✓ Loss = **categorical\_crossentropy** for multi-class
- ✓ Tracks: training & validation accuracy/loss

**history** stores accuracy/loss per epoch.

## e. Evaluate

```
python
```

```
test_loss, test_acc = model.evaluate(x_test, y_test)
```

- ✓ Gives final test accuracy & loss.

## Predict one example

```
python
```

```
plt.imshow(x_test[n])
predicted_value = model.predict(x_test)
print(np.argmax(y_test[n]))
print(np.argmax(predicted_value[n]))
```

- ✓ Displays the image
- ✓ Shows actual label & predicted label

## Plot training graphs

python

```
plt.plot(history['accuracy'])
plt.plot(history['val_accuracy'])
```

## ✓ Accuracy graph

python

```
plt.plot(history['loss'])
plt.plot(history['val_loss'])
```

## ✓ Loss graph

# ✓ PART 2 — Important Keywords, Libraries, Terms

## ✓ Important Libraries

Library	Purpose
<b>TensorFlow / Keras</b>	Neural network creation & training
<b>ImageDataGenerator</b>	Loads images from folders & preprocesses
<b>Matplotlib</b>	Plots graphs
<b>NumPy</b>	Matrix operations for input/output

## ✓ Important Keywords / Functions

Keyword	Meaning
<b>Sequential()</b>	Layer-by-layer simple model
<b>Flatten</b>	Converts image → 1D vector
<b>Dense layer</b>	Fully connected layer
<b>ReLU</b>	Activation function (non-linear)
<b>Softmax</b>	Converts outputs → probabilities
<b>SGD</b>	Stochastic Gradient Descent optimizer
<b>loss='categorical_crossentropy'</b>	Used for multi-class classification
<b>metrics=['accuracy']</b>	Tracks accuracy
<b>fit()</b>	Trains the model
<b>evaluate()</b>	Tests the model
<b>predict()</b>	Predicts output for inputs
<b>flow_from_directory()</b>	Loads images from folder structure
<b>rescale</b>	Normalizes pixel values

## ✓ PART 3 — Possible Viva Questions (SPPU Deep Learning Practical Exam)

### ✓ Theory-Based Questions

1. What is a feed-forward neural network?
  2. Why do we use the **Flatten** layer?
  3. Difference between **SGD** and **Adam** optimizer.
  4. What is the role of the **activation function**?
  5. Why do we use **ReLU** in hidden layers?
  6. What is **softmax** activation used for?
  7. Why do we normalize images?
  8. What is the difference between training loss and validation loss?
  9. What is overfitting? How do you detect it from graphs?
  10. What is categorical cross-entropy?
- 

### ✓ Program / Practical Questions

1. What will happen if we remove the Flatten layer?
  2. Why is the output layer size = 10 in MNIST?
  3. What happens if batch size is very large?
  4. Why do we use grayscale mode for MNIST?
  5. How does ImageDataGenerator help during training?
  6. What is one-hot encoding and why is it needed?
  7. What changes are required if using CIFAR-10 instead of MNIST?
  8. What happens if we change activation to sigmoid?
  9. Why do we use `.flow_from_directory()`?
  10. How will you calculate accuracy manually?
- 

### ✓ Advanced Viva Questions (If examiner goes deeper)

1. What is the difference between neural network and perceptron?
2. What is vanishing gradient? Does ReLU solve it?
3. What happens during backpropagation?
4. Why is softmax differentiable?
5. Explain the weight update rule in SGD.

If you want, I can also give you:

- A short **oral-viva-friendly summary (2 minutes)**
- A **printed PDF answer sheet** for submitting practical
- A **diagram of the network architecture**

Would you like those?



## 2. 1. what is mnist

2.x\_train, y\_train = train\_generator[0]

x\_test, y\_test = test\_generator[0] why [0]

3.what is this why isthis ? ✓ Flatten → converts 28×28 image → 784-dim vector

✓ Dense(128) → 128 neurons, activation = ReLU

✓ Dense(10) → 10 output classes, softmax gives probabilities

This is a Feed-Forward Neural Network (no convolution).

## 4.what is sgd

output related doubt also give me

Below are **simple + viva-perfect answers** for each question (SPPU Deep Learning practical style).

## ✓ 1. What is MNIST?

**MNIST = Modified National Institute of Standards and Technology dataset**

It is a dataset of:

- **70,000 handwritten digit images (0-9)**
- Image size: **28 × 28 pixels**
- **Grayscale (1 channel)**
- 60,000 training images
- 10,000 testing images

Used for:

- Image classification
- Testing neural networks
- Benchmark datasets in deep learning

👉 It is the “Hello World” dataset of deep learning.

## ✓ 2. Why do we use [0] here?

python

```
x_train, y_train = train_generator[0]
x_test, y_test = test_generator[0]
```

## ✓ Reason:

You used:

```
python
batch_size = 10000
```

Meaning the data generator will load **all 10,000 training images at once.**

The generator works like a list of batches:

- `train_generator[0]` = **first batch (contains all images)**
- `train_generator[1]` = second batch (in your case doesn't exist)

So `[0]` means:

✓ Give me the **first batch**, which contains  
 → `x_train (images)`  
 → `y_train (labels)`

Same for test data.

## ✓ 3. Explanation of: Flatten, Dense(128), Dense(10)

(Viva-ready simple explanation)

### ✓ Flatten Layer

```
arduino
```

Flatten converts input `image (28x28x1)` into a 1D vector of size 784.

Because a Dense layer accepts only **1D inputs**.

### ✓ Dense(128)

```
java
```

```
Dense layer = Fully Connected Layer
128 neurons
activation='relu'
```

Why ReLU?

- ✓ Faster training
- ✓ Reduces vanishing gradient
- ✓ Best for hidden layers

### ✓ Dense(10)

csharp

```
10 neurons = 10 output classes (0 to 9 digits)
activation='softmax'
```

Why Softmax?

- ✓ Converts raw values → probabilities
- ✓ Makes total sum = 1
- ✓ Used for **multi-class classification**

## ✓ Why this is a Feed Forward Network?

Because:

- ✓ Data flows **in one direction only**

(Input → Hidden → Output)

- ✓ No convolution layers
- ✓ No pooling
- ✓ No recurrence (no LSTM, GRU)

It is a **simple ANN / MLP (Multi-Layer Perceptron)**.

## ✓ 4. What is SGD?

**SGD = Stochastic Gradient Descent**

It is an **optimizer** used to update the weights during training.

### ✓ How it works (simple):

1. Takes one batch of data
2. Calculates loss
3. Computes gradients (slope)
4. Updates weights in opposite direction of gradient
5. Repeats every batch

### ✓ Why called “Stochastic”?

Because it updates using **random small batches**, not the full dataset.

### ✓ Advantages:

- Fast
- Less memory usage

- Works well for simple tasks like MNIST

## ✓ Formula (for viva):

```
ini
W_new = W_old - learning_rate * gradient
```

## ✓ If your doubt is about output:

When you do:

```
python
predicted_value = model.predict(x_test)
```

## ✓ What is predicted\_value[n]?

It is a 10-dimensional probability vector like:

```
csharp
[0.01, 0.03, 0.92, 0.01, 0.02, 0.00, 0.00, 0.01, 0.00, 0.00]
```

## ✓ What does np.argmax() do?

It finds the index of highest probability.

Example:

```
np.argmax([0.1, 0.7, 0.2]) = 1
```

So:

```
python
np.argmax(predicted_value[n])
```

= predicted digit by the model.

## ✓ Do you want a compact one-page viva answer sheet for printing?

I can prepare it if you want.



### 3. What is a feed-forward neural network?

Why do we use the Flatten layer?

Difference between SGD and Adam optimizer.

What is the role of the activation function?

Why do we use ReLU in hidden layers?

What is softmax activation used for?

Why do we normalize images?

What is the difference between training loss and validation loss?

What is overfitting? How do you detect it from graphs?

What is categorical cross-entropy?

#### Program / Practical Questions

What will happen if we remove the Flatten layer?

Why is the output layer size = 10 in MNIST?

What happens if batch size is very large?

Why do we use grayscale mode for MNIST?

How does ImageDataGenerator help during training?

What is one-hot encoding and why is it needed?

What changes are required if using CIFAR-10 instead of MNIST?

What happens if we change activation to sigmoid?

Why do we use .flow\_from\_directory()?

How will you calculate accuracy manually?

#### Advanced Viva Questions (If examiner goes deeper)

What is the difference between neural network and perceptron?

What is vanishing gradient? Does ReLU solve it?

What happens during backpropagation?

Why is softmax differentiable?

Explain the weight update rule in SGD.

for this questions give me 1-2 sentence answers

Below are **1-2 sentence, viva-perfect answers** for every question. Short, crisp, and examiner-friendly



## **THEORY QUESTIONS (Short Answers)**

### **1. What is a feed-forward neural network?**

A feed-forward neural network is a model where information flows only in one direction—input → hidden layers → output—without any loops or recurrence.

### **2. Why do we use the Flatten layer?**

Flatten converts a 2D/3D image into a 1D vector so that it can be passed into Dense (fully connected) layers.

### **3. Difference between SGD and Adam optimizer.**

SGD updates weights using a fixed learning rate, while Adam adapts the learning rate automatically using momentum and RMS propagation, usually training faster.

### **4. What is the role of the activation function?**

Activation functions introduce non-linearity, enabling the network to learn complex patterns beyond simple linear relationships.

### **5. Why do we use ReLU in hidden layers?**

ReLU is fast, avoids the vanishing gradient problem, and works better for deep networks.

### **6. What is softmax activation used for?**

Softmax converts outputs into probability scores that sum to 1, making it suitable for multi-class classification.

## 7. Why do we normalize images?

Normalization scales pixel values ( $0\text{--}255 \rightarrow 0\text{--}1$ ), helping the network train faster and with more stability.

## 8. What is the difference between training loss and validation loss?

Training loss is calculated on the training data, while validation loss is calculated on unseen data to measure generalization.

## 9. What is overfitting? How do you detect it from graphs?

Overfitting happens when the model memorizes training data but performs poorly on new data; graphs show training accuracy increasing while validation accuracy drops.

## 10. What is categorical cross-entropy?

It is a loss function used for multi-class classification that measures the difference between the true label distribution and predicted probabilities.

## PRACTICAL / PROGRAM QUESTIONS (Short Answers)

### 11. What will happen if we remove the Flatten layer?

The Dense layer will not accept image inputs because it needs a 1D vector, causing a shape mismatch error.

### 12. Why is the output layer size = 10 in MNIST?

MNIST has 10 classes (digits 0 to 9), so the network needs 10 output neurons.

### 13. What happens if batch size is very large?

Training becomes slower per epoch and may require more memory, and the model may generalize poorly.

### 14. Why do we use grayscale mode for MNIST?

MNIST images are originally single-channel grayscale, so using grayscale saves memory and preserves the dataset format.

### 15. How does ImageDataGenerator help during training?

It loads images efficiently, resizes them automatically, normalizes them, and can apply augmentation.

### 16. What is one-hot encoding and why is it needed?

One-hot encoding converts class numbers into binary vectors, which is required for categorical crossEntropy loss.

## 17. What changes are required if using CIFAR-10 instead of MNIST?

Image size becomes  $32 \times 32$ , color mode becomes RGB, and input shape changes from  $(28, 28, 1)$  to  $(32, 32, 3)$ .

## 18. What happens if we change activation to sigmoid?

Training becomes slower, gradients vanish more easily, and accuracy usually drops.

## 19. Why do we use `.flow_from_directory()`?

It automatically reads images from folders, labels them by directory name, resizes, batches, and normalizes them.

## 20. How will you calculate accuracy manually?

Accuracy = (Number of correct predictions / Total predictions)  $\times 100$ .

## ADVANCED QUESTIONS (Short Answers)

### 21. What is the difference between neural network and perceptron?

A perceptron is a single-layer classifier, while a neural network contains multiple layers and can learn complex non-linear functions.

### 22. What is vanishing gradient? Does ReLU solve it?

Vanishing gradient occurs when gradients become too small to update weights; ReLU reduces this problem because its gradient is constant for positive values.

### 23. What happens during backpropagation?

Errors are propagated backward through the network, and gradients are computed to update weights.

### 24. Why is softmax differentiable?

Softmax is a smooth function with a well-defined gradient, making it compatible with gradient descent.

### 25. Explain the weight update rule in SGD.

Weights are updated as:

**$W = W - \text{learning\_rate} \times \text{gradient}$ ,**

moving weights in the direction that reduces loss.

If you want, I can prepare:

- A **one-page viva sheet (PDF)**
- A **diagram of the network**
- A **super-short 2-minute oral explanation**

Just tell me!



#### 4. in code where is sgd used and how that formula is used

Here is the exact, simple explanation of **where SGD is used in your code** and **how the mathematical formula is applied inside TensorFlow**.

## ✓ 1. Where is SGD used in your code?

Right here 

```
python

model.compile(optimizer='sgd',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

### ✓ optimizer='sgd'

This tells TensorFlow:

**"Use Stochastic Gradient Descent to update the weights after every batch."**

After this line, during training:

```
python

history = model.fit(x_train, y_train, ...)
```

### → SGD starts working internally

- It updates weights after each batch
- Using the SGD formula

You never write the formula manually; TensorFlow applies it automatically.

## ✓ 2. How is the SGD formula used inside your model?

The core formula of SGD is:

```
ini

W_new = W_old - learning_rate * gradient
```

Where:

- **W\_old** = existing weight
- **gradient** = slope of loss wrt that weight

- **learning\_rate** = step size (default 0.01 if not specified)

## ✓ 3. How this formula works *in your code* (step-by-step during training)

### ✓ Step 1: Forward pass

The model does:

```
ini
y_pred = model(x_train)
```

### ✓ Step 2: Compute loss

Because you used:

```
python
loss='categorical_crossentropy'
```

TensorFlow calculates cross entropy between predicted and true labels.

### ✓ Step 3: Backpropagation

TensorFlow computes **gradients of loss** wrt each weight:

```
bash
dLoss/dW
```

### ✓ Step 4: Apply SGD formula

For **each neuron weight**, SGD updates it:

```
ini
W_new = W_old - lr * (dLoss/dW)
```

This is done for:

- weights in Dense(128)
- weights in Dense(10)

### ✓ Step 5: Repeat

SGD repeats this **every batch for every epoch**.

## ✓ 4. Important: You do NOT see the formula in code

Because TensorFlow handles everything internally.

You only specify:

```
python
optimizer='sgd'
```

And TensorFlow does:

- gradient calculation
- weight updates
- iteration
- minimization

Automatically.

## ✓ 5. Want to see SGD happening?

If you print model weights before and after training:

```
python
print(model.layers[1].get_weights()[0]) # weights before
```

After training:

```
python
print(model.layers[1].get_weights()[0]) # weights updated by SGD
```

You'll see values change → because SGD applied the formula.

## ✓ Summary for Viva (Best 2-line answer)

**SGD is used inside `model.compile()` as the optimizer.**

**TensorFlow applies the SGD update formula  $w = w - lr \times \text{gradient}$  internally during each batch of training.**

If you want, I can also give:

✓ Diagram of how SGD updates weights

- Simple numerical example of one SGD weight update
- Viva short 30-second explanation of SGD

Just tell me!



## 5. what is vanishing gradient problem

Here is the **shortest and clearest explanation** of the **vanishing gradient problem** for viva.

## Vanishing Gradient Problem (1-2 Sentence Explanation)

The **vanishing gradient problem** occurs when gradients become extremely small during backpropagation, causing earlier layers of a neural network to learn very slowly or not at all. This usually happens in deep networks using activation functions like **sigmoid** or **tanh**, where gradients shrink toward zero as values propagate backward.

## Slightly detailed explanation (if examiner asks "explain")

- In deep networks, errors are propagated backward using gradients.
- Certain activation functions (sigmoid/tanh) have derivatives between **0 and 1**, so multiplying these repeatedly across many layers makes gradients shrink.
- When gradients become very close to **zero**, earlier layers stop updating.
- As a result, the model fails to learn deep patterns.

## Symptoms of vanishing gradient

- Training accuracy improves very slowly
- Model gets stuck and doesn't learn
- Earlier layers have almost no weight updates

## How ReLU helps

ReLU has a gradient of **1** for positive values, so it does not shrink during backpropagation, helping avoid vanishing gradients.

If you want, I can also explain:

- Difference between *vanishing* and *exploding* gradients
- A simple graphical explanation
- A viva-ready 20-second answer for exam

Just tell me!



6. if practical is about feed forwards neural network then how can we back propagate

Great question — and this is **exactly the kind of viva trap** examiners ask in SPPU practicals.

Here is the **simple, perfect answer**:

## **If the practical is Feed-Forward Neural Network, how can we do Backpropagation?**

Because **backpropagation is not a type of neural network** —  
it is the **training algorithm** used to update weights **in any feed-forward network**.

A feed-forward network = architecture

Backpropagation = learning method

So **feed-forward refers to how data flows**

and **backpropagation refers to how errors flow backward to update weights**.

## **1-2 sentence viva answer**

"A feed-forward neural network passes data forward from input → output.

Backpropagation is the algorithm used to train it by sending the error backward and updating weights.  
So even though the network is feed-forward, training still uses backpropagation internally."

## **Simple explanation (20 seconds)**

- Feed-forward = structure (no loops, no backward data flow during prediction).
- Backpropagation = during training only, errors move backward to adjust weights.
- So prediction is forward; learning happens backward.

## **Why both exist in the same model?**

### **Forward pass:**

Compute output predictions.

### **Backward pass (backpropagation):**

Compute gradients → update weights.

This happens inside:

```
python
```

```
model.fit(...)
```

TensorFlow automatically performs **forward + backward** steps every batch.

## ✓ Best analogy for viva

Think of a feed-forward network like a calculator:

- **Forward:** You enter numbers → get result.
- **Backward:** If result is wrong, you adjust the calculator's internal knobs (weights).

The adjusting part = **backpropagation**.

## ✓ If you want, I can also give:

- ✓ A diagram showing forward + backward pass
- ✓ A 2-minute viva explanation
- ✓ How backprop works inside Dense layers

Just tell me!



7. so to summarize this just tell me how should i explain my code practical to external that what i my code about

Here is a **perfect, short, clear, viva-ready explanation** of your practical in **less than 1 minute**. Say this to the external examiner — it will cover everything.

## ✓ Viva-Ready Summary of Your Practical (Feed-Forward Neural Network using MNIST)

"In this practical, I have implemented a Feed-Forward Neural Network using TensorFlow and Keras to classify handwritten digits from the MNIST dataset. First, I loaded the MNIST images from my local folder using `ImageDataGenerator`, which also resized and normalized the images. Then I extracted the first batch to get `x_train`, `y_train` and `x_test`, `y_test`."

"My model is a simple ANN with: a Flatten layer to convert  $28 \times 28$  images into a 784-dim vector, one Dense hidden layer with 128 neurons using ReLU activation, and an output layer with 10 neurons and Softmax activation to classify digits 0–9."

"For training, I used the SGD optimizer and categorical cross-entropy loss. During `model.fit()`, TensorFlow automatically performs forward propagation to compute predictions and backpropagation to adjust weights. After training, I evaluated the model using test data and plotted training/validation accuracy and loss curves."

"Finally, I predicted an example digit to show that the model correctly identifies handwritten numbers."

## ✓ If you want it even shorter (30 seconds)

"My code builds a feed-forward neural network to classify MNIST digits. I loaded images with `ImageDataGenerator`, normalized them, and created a model with `Flatten` → `Dense(128, ReLU)` → `Dense(10, Softmax)`. I trained using SGD, evaluated accuracy, and plotted training and validation curves. The model finally predicts digits correctly using softmax output."

## ✓ If you want a 15-second version (very short)

"This practical trains a simple feed-forward neural network on MNIST. I load and normalize images, flatten them, pass through a Dense ReLU layer and a Softmax output layer, train using SGD, evaluate accuracy, and plot training results."

If you want, I can also give you:

- A full **ready-to-print practical write-up**
- A **diagram** of your model architecture
- A **viva cheat-sheet** with Q&A

Just tell me!



## 8. and what does sgd do

Here is the **shortest, clearest, viva-perfect explanation of what SGD does:**

### **What does SGD do? (1-2 sentences)**

**SGD (Stochastic Gradient Descent) updates the weights of the neural network during training so that the loss decreases. It calculates the gradient of the loss and moves the weights slightly in the direction that reduces error.**

### **Simple explanation (20 seconds)**

SGD looks at a small batch of data, measures how wrong the model is (loss), computes the slope (gradient), and then adjusts the weights using:

```
ini
W_new = W_old - learning_rate * gradient
```

This process is repeated for every batch and every epoch until the model learns.

### **Even simpler (examiner-friendly)**

**"SGD trains the model by slowly correcting weights after every batch so the model predicts more accurately."**

If you want, I can also give:

- A tiny numerical example of SGD updating one weight
- A 5-second answer for viva
- A diagram showing SGD steps



## 9. ok now #a. Import the necessary packages

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.layers import Input, Conv2D, Dense, Flatten, Dropout
from tensorflow.keras.layers import GlobalMaxPooling2D, MaxPooling2D
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Model
from keras.utils import to_categorical
import matplotlib.pyplot as plt
import numpy as np
```

## #b. Load the training and testing data (MNIST/CIFAR10)

```
train_data_dir = 'cifar-10-img/train'
test_data_dir = 'cifar-10-img/test'

# Image data generator for training data
train_datagen = ImageDataGenerator(
    rescale=1.0/255
)

# Image data generator for testing data
test_datagen = ImageDataGenerator(
    rescale=1.0/255
)

# Create data generators
train_batch_size = 20000
train_generator = train_datagen.flow_from_directory(
    train_data_dir,
    target_size=(32, 32), # Resize images to 32x32
    batch_size=train_batch_size,
    class_mode='sparse',
    shuffle=True,
)

# Load test data without labels (class_mode=None)
test_batch_size = 1000
test_generator = test_datagen.flow_from_directory(
    test_data_dir,
```

```
target_size=(32, 32), # Resize images to 32x32  
batch_size=test_batch_size,  
class_mode='sparse',  
shuffle=True,  
)
```

```
x_train, y_train = train_generator[0]  
x_test, y_test = test_generator[0]
```

```
y_train = y_train.flatten()  
y_test = y_test.flatten()
```

### #c. Define the network architecture using Keras

```
# number of classes  
K = len(set(y_train))  
# calculate total number of classes  
# for output layer  
print("number of classes:", K)  
i = Input(shape=x_train[0].shape)  
x = Conv2D(32, (3, 3), activation='relu', padding='same')(i)  
x = BatchNormalization()(x)  
x = Conv2D(32, (3, 3), activation='relu', padding='same')(x)  
x = BatchNormalization()(x)  
x = MaxPooling2D((2, 2))(x)  
  
x = Conv2D(64, (3, 3), activation='relu', padding='same')(x)  
x = BatchNormalization()(x)  
x = Conv2D(64, (3, 3), activation='relu', padding='same')(x)  
x = BatchNormalization()(x)  
x = MaxPooling2D((2, 2))(x)
```

```
x = Conv2D(128, (3, 3), activation='relu', padding='same')(x)
x = BatchNormalization()(x)
x = Conv2D(128, (3, 3), activation='relu', padding='same')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2, 2))(x)

x = Flatten()(x)
x = Dropout(0.2)(x)

# Hidden layer
x = Dense(1024, activation='relu')(x)
x = Dropout(0.2)(x)

# last hidden layer i.e.. output layer
x = Dense(K, activation='softmax')(x)

model = Model(i, x)

#D.Train the model using SGD
model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
history = model.fit(x_train, y_train, epochs=10, validation_data=(x_test, y_test))

#e. Evaluate the network
test_loss, test_acc = model.evaluate(x_test, y_test)
print("Loss: ", test_loss)
print("Accuracy: ", test_acc)

import matplotlib.pyplot as plt
y = to_categorical(y_test)
predicted_value = model.predict(x_test)
```

```
n = 10
plt.imshow(x_test[n])
print("Actual Number: ",labels[np.argmax(y[n])])
print("Predicted Number: ", labels[np.argmax(predicted_value[n])])
```

```
#f. Plot the training loss and accuracy
```

```
history = history.history
history.keys()
```

```
plt.plot(history['accuracy'])
plt.plot(history['val_accuracy'])
plt.title("Model Accuracy")
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['Train', "Validation"], loc='upper left')
```

```
plt.plot(history['loss'])
plt.plot(history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
```

```
plt.plot(history['accuracy'])
plt.plot(history['val_accuracy'])
plt.plot(history['loss'])
plt.plot(history['val_loss'])
plt.title('Training Loss and accuracy')
plt.ylabel('accuracy/Loss')
plt.xlabel('epoch')
plt.legend(['accuracy', 'val_accuracy', 'loss', 'val_loss'])
```

do the same for this

Below is the **perfect viva-ready summary** of your **CIFAR-10 CNN practical**, just like we did for the MNIST feed-forward NN.

This is the **1-minute explanation you give to the external examiner**, along with **30-sec** and **15-sec** versions.

## ✓ Viva-Ready Summary of Your CNN Practical (1 Minute)

**"In this practical, I implemented a Convolutional Neural Network (CNN) using TensorFlow/Keras to classify CIFAR-10 images. I first loaded the CIFAR-10 dataset from local folders using ImageDataGenerator, which resizes images to 32×32, normalizes them (rescale=1/255), and loads them in batches."**

**"I extracted the first batch using train\_generator[0] and test\_generator[0] which gave me x\_train, y\_train, x\_test, and y\_test. Since CIFAR-10 has 10 classes, I kept the labels in sparse form and used sparse\_categorical\_crossentropy as the loss function."**

**"My CNN architecture has multiple blocks of Conv2D → BatchNormalization → MaxPooling layers. These layers extract spatial features from RGB images. After convolution blocks, I flattened the**

features, added Dropout for regularization, then used a Dense(1024, ReLU) hidden layer, and finally a Dense(K, Softmax) output layer for 10-class classification.”

“I compiled the model using the SGD optimizer and trained it for 10 epochs. During training, TensorFlow automatically uses backpropagation to update CNN filters. I evaluated accuracy on the test set, predicted sample images, and plotted training vs. validation accuracy and loss curves.”

## Shorter (30-second) version

“This practical builds a CNN for CIFAR-10 classification. Images are loaded using `ImageDataGenerator`, resized to 32×32, and normalized. My model uses stacked Conv2D layers with BatchNormalization and MaxPooling for feature extraction, followed by Flatten, Dropout, Dense(1024), and a Softmax output layer. I trained using SGD with sparse cross-entropy, evaluated accuracy, tested predictions, and plotted loss/accuracy graphs.”

## Very Short (15-second) version

“I implemented a CNN on CIFAR-10 using Conv2D, BatchNorm, MaxPooling, Dense, and Softmax layers. Images are loaded with `ImageDataGenerator`, normalized, trained using SGD, evaluated on the test set, and training graphs are plotted.”

## If the External Asks “Explain Your Code Part-by-Part”

Use this breakdown:

### Part A – Import Packages (Explain in 10 seconds)

- `tensorflow/keras` → Build and train CNN
- `layers` → Conv2D, Dense, Flatten, etc.
- `ImageDataGenerator` → Load images from folders
- `matplotlib` → Plot graphs
- `numpy` → Numerical operations

### Part B – Load CIFAR-10 Data

-  `ImageDataGenerator(rescale=1/255)` → Normalize pixels to improve training
-  `flow_from_directory()` → Reads images from folder structure
-  `target_size=(32,32)` → Resize to CIFAR-10 size
-  `class_mode='sparse'` → Labels stored as integers (0-9)

```
python
```

```
x_train, y_train = train_generator[0]
x_test, y_test = test_generator[0]
```

✓ [0] gives **first batch**, which contains all images because batch size is 20,000.

✓ `y_train.flatten()` → Converts label shape from (20000,1) to (20000,)

## ✓ Part C – CNN Model Architecture

### Block-1

- Conv2D(32) → learn edges
- BatchNorm → stabilize learning
- Conv2D(32)
- BatchNorm
- MaxPooling → reduce size

### Block-2

- Conv2D(64)
- BatchNorm
- Conv2D(64)
- BatchNorm
- MaxPooling

### Block-3

- Conv2D(128)
- BatchNorm
- Conv2D(128)
- BatchNorm
- MaxPooling

### Fully Connected Layers

- Flatten → convert 3D → 1D
- Dropout(0.2) → reduce overfitting
- Dense(1024, relu) → learning high-level features
- Dropout

- Final Dense(K, softmax) → K = 10 classes

## Softmax → probability output for 10 classes

## Part D – Training with SGD

python

```
model.compile(optimizer='sgd',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

- SGD updates weights using:  
 **$W_{\text{new}} = W_{\text{old}} - lr \times \text{gradient}$**
- Loss = sparse cross-entropy → works with integer labels
- Backpropagation automatically handled by TensorFlow

## Part E – Evaluation

python

```
test_loss, test_acc = model.evaluate(x_test, y_test)
```

 Gives model performance on unseen data.

## Part F – Prediction + Plotting

- Used `model.predict()` to classify images
- Used `plt.imshow()` to display sample image
- Plotted accuracy and loss graphs for training vs validation  
→ To detect overfitting/underfitting

If you want, I can also provide:

-  Viva Q&A for CNN (Conv2D, BatchNorm, Pooling, Dropout, Softmax)
-  A written **practical journal ready PDF**
-  A model architecture diagram

Just tell me!



10. explain me this code in three parts for viva first : function by function which does what , important keywords and libraries and third possible questions asked maybe theory based or practical based sppu deep learnig context

Below is the **perfect 3-part viva explanation** for your **CIFAR-10 CNN practical**, exactly like you asked:

- PART 1 — Function-by-function explanation (your entire code, block by block)**
- PART 2 — Important keywords, layers, libraries used**
- PART 3 — Viva questions (theory + practical, SPPU-style)**

This is everything you need for a full external viva.

## **PART 1 — FUNCTION-BY-FUNCTION EXPLANATION (CODE WALKTHROUGH)**

### **A. Import Necessary Packages**

python

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.layers import Input, Conv2D, Dense, Flatten, Dropout
from tensorflow.keras.layers import GlobalMaxPooling2D, MaxPooling2D
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Model
from keras.utils import to_categorical
import matplotlib.pyplot as plt
import numpy as np
```

- TensorFlow/Keras** → Build neural network
- Conv2D, Dense, Flatten** → CNN layers
- BatchNormalization** → Stabilizes training
- ImageDataGenerator** → Loads images from folder
- Model API** → Functional model creation
- Matplotlib** → Plot accuracy + loss
- NumPy** → Numerical processing

### **B. Load Training & Testing Data (CIFAR-10)**

#### **Folder paths**

```
python
```

```
train_data_dir = 'cifar-10-img/train'  
test_data_dir = 'cifar-10-img/test'
```

## Image Generators

```
python
```

```
train_datagen = ImageDataGenerator(rescale=1.0/255)  
test_datagen = ImageDataGenerator(rescale=1.0/255)
```

- Normalizes pixel values (0–255 → 0–1)
- Helps faster & stable training

## Load training images

```
python
```

```
train_generator = train_datagen.flow_from_directory(  
    train_data_dir,  
    target_size=(32, 32),  
    batch_size=20000,  
    class_mode='sparse',  
    shuffle=True)
```

- Reads all images as RGB (32×32×3)
- `class_mode='sparse'` → labels stored as integers (0–9)
- Large batch loads **all training data in one go**

## Load test images

```
python
```

```
test_generator = test_datagen.flow_from_directory(  
    test_data_dir,  
    target_size=(32, 32),  
    batch_size=1000,  
    class_mode='sparse',  
    shuffle=True)
```

- Same preprocessing as training data

## Extract first batch

```
python
```

```
x_train, y_train = train_generator[0]  
x_test, y_test = test_generator[0]
```

[0] gives **the first batch**, which contains **all images** because batch size is large

## Flatten labels

```
python
```

```
y_train = y_train.flatten()
y_test = y_test.flatten()
```

Turns shape from  $(20000, 1) \rightarrow (20000,)$

Required for sparse cross-entropy.

## C. Define CNN Architecture (Functional API)

### Number of classes

```
python
```

```
K = len(set(y_train))
```

CIFAR-10  $\rightarrow K = 10$

### Input layer

```
python
```

```
i = Input(shape=x_train[0].shape)
```

Input shape =  $(32, 32, 3)$

### Convolution Block 1

```
python
```

```
x = Conv2D(32, (3,3), activation='relu', padding='same')(i)
x = BatchNormalization()(x)
x = Conv2D(32, (3,3), activation='relu', padding='same')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2,2))(x)
```

Extracts low-level features

BatchNorm speeds up training

MaxPooling reduces size (downsampling)

### Convolution Block 2 (same pattern, more filters)

Learns deeper features like textures

## Convolution Block 3 (even deeper)

Learns complex shapes & patterns

## Flatten + Regularization

```
python
```

```
x = Flatten()(x)
x = Dropout(0.2)(x)
```

Flatten → Convert 3D feature map to 1D

Dropout → Prevents overfitting

## Fully Connected Layers

```
python
```

```
x = Dense(1024, activation='relu')(x)
x = Dropout(0.2)(x)
```

High-level feature learning

More dropout for regularization

## Softmax Output Layer

```
python
```

```
x = Dense(K, activation='softmax')(x)
model = Model(i, x)
```

10 neurons → 10 classes

Softmax outputs probabilities

## D. Train the Model (SGD)

```
python
```

```
model.compile(optimizer='sgd',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

- SGD** updates weights with gradient descent
- Sparse cross-entropy** works with integer labels
- Tracks accuracy

```
python
```

```
history = model.fit(x_train, y_train,
                     epochs=10,
                     validation_data=(x_test, y_test))
```

- Training + validation in each epoch
- TensorFlow automatically uses **backpropagation**

## E. Evaluate the Model

```
python
```

```
test_loss, test_acc = model.evaluate(x_test, y_test)
```

- Shows final test accuracy & loss

## F. Predict one image

```
python
```

```
predicted_value = model.predict(x_test)
plt.imshow(x_test[n])
```

- Displays sample test image
- Compares actual vs predicted class

## G. Plot Accuracy & Loss

- Plots:

- Training accuracy
- Validation accuracy
- Training loss
- Validation loss

These graphs help detect overfitting.

## ✓ PART 2 — IMPORTANT KEYWORDS, LAYERS & LIBRARIES

### ✓ Important Libraries

Library	Purpose
TensorFlow/Keras	Build CNN
ImageDataGenerator	Load images from folders
Conv2D	Extract features
MaxPooling2D	Downsample
BatchNormalization	Stabilize/accelerate training
Dropout	Prevent overfitting
Matplotlib	Plot graphs
NumPy	Numerical calculations

### ✓ Important Keywords

Keyword	Meaning
<code>Input()</code>	Defines input shape
<code>Conv2D()</code>	Convolutional layer
<code>padding='same'</code>	Output size preserved
<code>ReLU</code>	Activation function
<code>BatchNormalization</code>	Normalizes activations
<code>MaxPooling2D</code>	Reduces width+height
<code>Flatten</code>	Converts to 1D
<code>Dense</code>	Fully connected layer
<code>Softmax</code>	Output probabilities
<code>SGD</code>	Optimizer
<code>sparse_categorical_crossentropy</code>	Loss for integer labels

## ✓ PART 3 — POSSIBLE VIVA QUESTIONS (SPPU)

### ✓ Theory Viva Questions

1. What is a CNN?
2. Why do we use Conv2D layers?
3. What is the purpose of BatchNormalization?

4. Why use MaxPooling?
  5. What is the difference between RGB and grayscale images?
  6. What is ReLU and why is it used?
  7. Why do we use softmax in the output layer?
  8. What is Dropout? Why do we use it?
  9. What is sparse categorical cross-entropy?
  10. What is overfitting and how do graphs show it?
- 

## Practical Viva Questions

1. Why did you resize to 32×32?
  2. Why use `class_mode='sparse'`?
  3. Why batch size = 20,000?
  4. Why use three Conv blocks?
  5. What happens if we remove Dropout?
  6. Why do we normalize images (rescale=1/255)?
  7. What does `model.compile()` do?
  8. What is the role of SGD?
  9. What is the effect of increasing epochs?
  10. Why do we use Flatten before Dense?
- 

## Advanced Viva Questions

1. Explain backpropagation in CNN.
  2. What is a feature map?
  3. What is filter/kernel size?
  4. What is padding?
  5. Why BatchNorm helps reduce internal covariate shift?
  6. Why deep CNN performs better than ANN for images?
- 

If you want, I can also give you:

-  A **one-page viva answer sheet**
-  A **diagram of your CNN architecture**
-  A **printable practical write-up PDF**

Just tell me!



## 11. viva questions with answers

Below are **SPPU-focused Viva Questions with short, perfect answers** for your **CIFAR-10 CNN Practical**, covering:

- Theory questions
- Practical/program questions
- Deep learning conceptual questions
- Advanced/external-questions

All answers are **1-2 sentences**, viva-ready.

## PART A — Theory Viva Questions (CNN + DL Basics)

### 1. What is a Convolutional Neural Network (CNN)?

A CNN is a deep learning model designed to automatically learn features from images using convolution filters and pooling layers.

### 2. Why do we use Conv2D layers?

Conv2D extracts patterns such as edges, textures, and shapes by sliding filters over the image.

### 3. What is ReLU activation?

ReLU =  $\max(0, x)$ . It removes negative values and helps avoid the vanishing gradient problem.

### 4. What is MaxPooling?

MaxPooling reduces the spatial size of feature maps by taking the maximum value, helping reduce computation and overfitting.

### 5. What is Batch Normalization?

BatchNorm normalizes activations during training, making training faster and more stable.

### 6. What is the Softmax activation used for?

Softmax converts the output of the final layer into probabilities that sum to 1, used for multi-class classification.

### 7. Why do we normalize images (1/255)?

To scale pixel values to 0-1, which speeds up training and prevents large input values from causing instability.

## 8. What is Dropout and why do we use it?

Dropout randomly disables neurons during training to prevent overfitting.

## 9. What is overfitting?

When a model performs well on training data but poorly on test data; seen when validation loss increases while training loss decreases.

## 10. What is the purpose of Flatten?

Flatten converts the final 3D feature map into a 1D vector for the Dense layer.

## PART B — Practical / Code-Based Viva Questions

### 11. Why is the input size (32, 32, 3) in CIFAR-10?

Because CIFAR-10 contains color images of size 32×32 with 3 channels (RGB).

### 12. Why do you use `class_mode='sparse'`?

Because labels are integers (0–9) and sparse cross-entropy works directly with integer labels.

### 13. What does `train_generator[0]` do?

It returns the first batch of images and labels; since batch size is large, it contains the entire dataset.

### 14. Why use multiple Conv layers before pooling?

Stacking Conv layers helps the network learn more complex and deeper features before reducing size with pooling.

### 15. Why is the output layer size = 10?

CIFAR-10 has 10 classes, so we need 10 neurons in the final Softmax layer.

### 16. Why do we use `sparse_categorical_crossentropy`?

Because labels are integers, not one-hot encoded vectors.

### 17. What does SGD do in this model?

SGD updates model weights using gradients to reduce loss using:

$$W_{\text{new}} = W_{\text{old}} - \text{learning\_rate} \times \text{gradient}$$

### 18. What happens if we remove Batch Normalization?

Training becomes slower, and the model may get stuck or become unstable.

## 19. What happens if we remove Dropout?

The model may overfit, meaning it learns training data too well but performs poorly on test data.

## 20. Why do you shuffle data?

Shuffling prevents the model from learning patterns based on data order and improves generalization.

# ✓ PART C — Deep Learning Concepts (Common Viva Questions)

## 21. What is backpropagation?

Backpropagation calculates gradients of loss with respect to weights and updates them to minimize error.

## 22. What is gradient descent?

An optimization method that updates weights by moving in the negative direction of the gradient to reduce loss.

## 23. What is a feature map?

The output of a convolution layer showing extracted features such as edges or textures.

## 24. What is a filter/kernel in CNN?

A small matrix (e.g.,  $3 \times 3$ ) that slides over the image to detect patterns.

## 25. What is padding='same'?

Padding keeps output size equal to input size by adding zeros around the borders.

## 26. Why use the Functional API instead of Sequential?

Functional API allows complex architectures with multiple inputs, outputs, or branching.

## 27. What is the vanishing gradient problem?

Gradients become too small during backpropagation, slowing or stopping learning in deep networks.

## 28. What is the role of activation functions?

They introduce non-linearity, enabling the network to learn complex patterns.

## 29. What is the difference between training and validation accuracy?

Training accuracy is measured on training data; validation accuracy on unseen data to test generalization.

### 30. Why use Softmax instead of Sigmoid for multi-class classification?

Softmax outputs a probability distribution over multiple classes; sigmoid is only for binary classification.

## ✓ PART D — Advanced Viva Questions (Asked by strict externals)

### 31. Why does CNN perform better than ANN for images?

CNNs preserve spatial features and reduce parameters using kernels, unlike ANN which flattens data and loses spatial information.

### 32. What is internal covariate shift?

Changes in layer input distribution during training; BatchNorm helps reduce it.

### 33. What is the role of the learning rate?

Controls how big the step is when updating weights; too high diverges, too low makes training slow.

### 34. Difference between SGD and Adam optimizer.

SGD uses fixed learning rate; Adam adapts learning rates automatically using momentum and RMS.

### 35. How does MaxPooling reduce overfitting?

By reducing feature map size, it removes unnecessary details and forces the network to learn only dominant features.

### 36. Why are deeper CNNs better?

Deeper networks capture more complex patterns and hierarchical features.

## ✓ PART E — Prediction / Output Viva Questions

### 37. What does `model.predict()` return?

It returns probability scores for each class.

### 38. What does `np.argmax()` do?

It returns the index of the highest probability class.

### 39. Why do we convert labels using `to_categorical()`?

To convert integer labels into one-hot vectors when training with categorical cross-entropy.

## 40. Why plot loss and accuracy curves?

To observe learning behavior, detect overfitting, and understand model convergence.

If you want, I can also prepare:

- A **PDF of all viva questions + your code summary**
- A **diagram** of the architecture
- A **1-minute scripted viva answer** for flawless presentation



## 12. What is a CNN?

Why do we use Conv2D layers?

What is the purpose of BatchNormalization?

Why use MaxPooling?

What is the difference between RGB and grayscale images?

What is ReLU and why is it used?

Why do we use softmax in the output layer?

What is Dropout? Why do we use it?

What is sparse categorical cross-entropy?

What is overfitting and how do graphs show it?

### Practical Viva Questions

Why did you resize to 32×32?

Why use class\_mode='sparse'?

Why batch size = 20,000?

Why use three Conv blocks?

What happens if we remove Dropout?

Why do we normalize images (rescale=1/255)?

What does model.compile() do?

What is the role of SGD?

What is the effect of increasing epochs?

Why do we use Flatten before Dense?

### Advanced Viva Questions

Explain backpropagation in CNN.

What is a feature map?

What is filter/kernel size?

What is padding?

Why BatchNorm helps reduce internal covariate shift?

Why deep CNN performs better than ANN for images?

If you want, I can als

Below are **short, crisp, viva-perfect 1-2 sentence answers** for every question you listed. These are exactly what externals expect in SPPU Deep Learning viva.

## ✓ THEORY VIVA QUESTIONS (CNN)

### 1. What is a CNN?

A CNN (Convolutional Neural Network) is a deep learning model designed to automatically extract spatial features from images using convolution and pooling operations.

### 2. Why do we use Conv2D layers?

Conv2D layers apply learnable filters over an image to detect features like edges, textures, shapes, and patterns.

### 3. What is the purpose of BatchNormalization?

BatchNorm normalizes layer outputs to stabilize and speed up training, and helps reduce internal covariate shift.

### 4. Why use MaxPooling?

MaxPooling reduces the spatial size of feature maps, removes noise, and helps prevent overfitting by keeping only the strongest features.

### 5. What is the difference between RGB and grayscale images?

RGB images have 3 color channels (Red, Green, Blue), while grayscale images have only 1 channel representing intensity.

### 6. What is ReLU and why is it used?

ReLU =  $\max(0, x)$ . It introduces non-linearity, trains faster, and helps avoid the vanishing gradient problem.

## 7. Why do we use softmax in the output layer?

Softmax converts raw outputs into probability values that sum to 1, used for multi-class classification.

## 8. What is Dropout? Why do we use it?

Dropout randomly disables neurons during training to prevent overfitting and improve generalization.

## 9. What is sparse categorical cross-entropy?

It is a loss function used when labels are integer-encoded instead of one-hot encoded for multi-class classification.

## 10. What is overfitting and how do graphs show it?

Overfitting happens when the model learns training data too well but performs poorly on test data; graphs show training accuracy increasing while validation accuracy decreases.

## PRACTICAL / CODE VIVA QUESTIONS

### 11. Why did you resize to 32×32?

Because CIFAR-10 images are originally 32×32 pixels, so resizing preserves their natural resolution.

### 12. Why use `class_mode='sparse'`?

Because labels are integers (0–9), and sparse cross-entropy directly accepts integer labels without one-hot encoding.

### 13. Why batch size = 20,000?

To load the entire training dataset in one batch since the images are small and CIFAR-10 is lightweight.

### 14. Why use three Conv blocks?

Each block learns deeper features — first edges, then textures, then shapes and object-level patterns.

### 15. What happens if we remove Dropout?

The model may overfit because all neurons stay active, increasing the risk of memorizing training data.

### 16. Why do we normalize images (`rescale=1/255`)?

Normalization converts pixel values from 0–255 to 0–1, ensuring faster and more stable training.

## 17. What does `model.compile()` do?

It sets the loss function, optimizer, and evaluation metrics so the model knows how to train.

## 18. What is the role of SGD?

SGD updates weights using gradients to minimize loss, applying the formula:

$$W_{\text{new}} = W_{\text{old}} - \text{learning\_rate} \times \text{gradient}$$

## 19. What is the effect of increasing epochs?

Increasing epochs usually improves learning, but too many epochs can lead to overfitting.

## 20. Why do we use Flatten before Dense?

Flatten converts 3D feature maps into a 1D vector so the Dense layer can process them.

## ADVANCED VIVA QUESTIONS

### 21. Explain backpropagation in CNN.

Backpropagation calculates gradients of loss with respect to each weight (filters and neurons) and updates them to reduce error.

### 22. What is a feature map?

A feature map is the output of a convolution layer that highlights detected features like edges or textures.

### 23. What is filter/kernel size?

It is the dimension of the sliding window (e.g., 3x3) used by Conv2D to extract features.

### 24. What is padding?

Padding adds zeros around the image borders to preserve spatial dimensions after convolution.

### 25. Why BatchNorm helps reduce internal covariate shift?

BatchNorm keeps layer inputs stable by normalizing them, preventing shifts in data distribution during training.

### 26. Why deep CNN performs better than ANN for images?

CNNs preserve spatial relationships using filters, while ANNs flatten data and lose spatial information.

If you want, I can make you a **one-page viva cheat sheet, diagram of your CNN architecture, or a printed practical journal version.**



13. #3Build the Image classification model by dividing the model into following 4 stages:
- Loading and preprocessing the image data
  - Defining the model's architecture
  - Training the model
  - Estimating the model's performance

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical
from tensorflow.keras import layers, models
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt
```

```
#a. Loading and preprocessing the image data
train_data_dir = 'mnist-jpg/train'
test_data_dir = 'mnist-jpg/test'
```

```
# Image data generator for training data
train_datagen = ImageDataGenerator(
    rescale=1.0/255
)
```

```
# Image data generator for testing data
test_datagen = ImageDataGenerator(
    rescale=1.0/255
)
```

```
# Create data generators
train_batch_size = 10000
train_generator = train_datagen.flow_from_directory(
    train_data_dir,
    target_size=(28, 28), # Resize images to 28x28
    batch_size=train_batch_size,
    class_mode='categorical',
    color_mode='grayscale',# Use 'categorical' for one-hot encoded labels
    shuffle=True,
)
```

```
# Load test data without labels (class_mode=None)
test_batch_size = 2000
test_generator = test_datagen.flow_from_directory(
    test_data_dir,
    target_size=(28, 28), # Resize images to 28x28
    batch_size=test_batch_size,
    class_mode='categorical', # Use 'categorical' for one-hot encoded labels
    color_mode='grayscale',
    shuffle=True,
)
```

```
x_train, y_train = train_generator[0]
x_test, y_test = test_generator[0]
```

```
print(x_train.shape, y_train.shape)
```

### #b. Defining the model's architecture

```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

### #Cc. Training the model

```
model.fit(x_train, y_train, epochs=5, batch_size=64, validation_data=(x_test, y_test))
```

### #d.d. Estimating the model's performance

```
test_loss, test_accuracy = model.evaluate(x_test, y_test)
print("Loss: ", test_loss)
print("Accuracy: ", test_accuracy)
```

```
n = 30
plt.imshow(x_test[n])
predicted_value = model.predict(x_test)
print("Actual Number: ", np.argmax(y_test[n]))
print("Predicted Number: ", np.argmax(predicted_value[n]))
```

explain me this code in three parts for viva first : function by function which does what , important keywords and libraries and third possible questions asked maybe theory based or practical based sppu deep learnig context

Below is the **perfect 3-part viva explanation** for your **MNIST CNN image-classification practical**, exactly in the same format used in SPPU Deep Learning exams:

- Part 1 – Function-by-function explanation (your whole code explained step-by-step)**
- Part 2 – Important keywords, layers, libraries**
- Part 3 – Viva questions with short answers (theory + practical)**

## PART 1 — Function-by-Function Explanation (VIVA)

### A. Loading & Preprocessing Image Data

#### Imports

```
python

import numpy as np
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical
from tensorflow.keras import layers, models
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt
```

- tensorflow/keras** → building and training CNN
- ImageDataGenerator** → load and preprocess images from folder
- matplotlib** → plot images
- numpy** → numerical computations

#### Loading dataset from folder

```
python

train_data_dir = 'mnist-jpg/train'
test_data_dir = 'mnist-jpg/test'
```

- Your MNIST dataset is stored in 10 folders (0–9).

#### Image Preprocessing

```
python
```

```
train_datagen = ImageDataGenerator(rescale=1.0/255)
test_datagen = ImageDataGenerator(rescale=1.0/255)
```

- ✓ Normalizes pixel values (0–255 → 0–1)
- ✓ Makes training stable & fast

## Create Data Generators

```
python
```

```
train_generator = train_datagen.flow_from_directory(
    train_data_dir,
    target_size=(28, 28),
    batch_size=10000,
    class_mode='categorical',
    color_mode='grayscale',
    shuffle=True)
```

- ✓ Reads all training images at once
- ✓ Converts them to grayscale (1 channel)
- ✓ One-hot encodes labels (**categorical**)
- ✓ Resizes images to 28×28

Same for test generator.

## Load first batch

```
python
```

```
x_train, y_train = train_generator[0]
x_test, y_test = test_generator[0]
```

- ✓ [0] selects the first batch (containing all images due to large batch size)

## Check shapes

```
python
```

```
print(x_train.shape, y_train.shape)
```

- ✓ Confirms data correctly loaded (10000, 28, 28, 1)
- ✓ Labels: (10000, 10)

## ✓ B. Defining the Model Architecture

```
python
```

```
model = models.Sequential()
```

- ✓ Simple stack of layers

## 1. Convolution Layer

```
python
```

```
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
```

- ✓ 32 filters → extracts features
- ✓ Kernel size 3×3
- ✓ ReLU for non-linearity
- ✓ Input = grayscale MNIST image

## 2. Pooling Layer

```
python
```

```
model.add(layers.MaxPooling2D((2, 2)))
```

- ✓ Reduces image size
- ✓ Keeps important features

## 3. Flatten Layer

```
python
```

```
model.add(layers.Flatten())
```

- ✓ Converts 2D feature maps → 1D vector to feed Dense layers

## 4. Dense Hidden Layer

```
python
```

```
model.add(layers.Dense(64, activation='relu'))
```

- ✓ Learns high-level patterns

## 5. Output Layer

```
python
```

```
model.add(layers.Dense(10, activation='softmax'))
```

- ✓ Softmax gives probability distribution
- ✓ 10 classes (digits 0–9)

## Compile Model

```
python
```

```
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

- ✓ Adam optimizer → fast learning
- ✓ Categorical cross-entropy → used for multi-class one-hot labels
- ✓ Track accuracy

## ✓ C. Training the Model

```
python
```

```
model.fit(x_train, y_train, epochs=5, batch_size=64,
           validation_data=(x_test, y_test))
```

- ✓ Trains model
- ✓ 5 epochs
- ✓ Batch size = 64
- ✓ Backpropagation updates weights automatically

## ✓ D. Evaluating Model Performance

```
python
```

```
test_loss, test_accuracy = model.evaluate(x_test, y_test)
```

- ✓ Measures model performance on unseen data

## ✓ Prediction Example

python

```
plt.imshow(x_test[n])
predicted_value = model.predict(x_test)
print(np.argmax(y_test[n]))
print(np.argmax(predicted_value[n]))
```

- ✓ Shows sample test image
- ✓ Displays actual & predicted digit

## ✓ PART 2 — Important Keywords / Libraries

### ✓ Important Libraries

Library	Purpose
TensorFlow/Keras	Neural network creation & training
ImageDataGenerator	Load and preprocess images
Conv2D	Convolution layers
MaxPooling2D	Downsampling
Dense	Fully connected layer
Flatten	Convert 2D to 1D
Adam	Optimizer
Matplotlib	Visualization

### ✓ Important Keywords

Keyword	Meaning
Convolution	Extracts image features
ReLU	Activation to avoid vanishing gradients
Softmax	Converts to probability distribution
Categorical cross-entropy	Loss for multi-class, one-hot labels
Epoch	One complete pass through the dataset
Batch size	Number of samples processed before updating weights
Backpropagation	Error sent backwards to update weights

## ✓ PART 3 — Viva Questions with Perfect Answers (SPPU)

### ✓ Theory Questions

## 1. What is MNIST?

A dataset of 70,000 handwritten digit images ( $28 \times 28$  grayscale).

## 2. What is a CNN?

A CNN is a deep learning model that learns image features using Convolution and Pooling.

## 3. Why use Conv2D?

Conv2D extracts spatial patterns such as edges and shapes.

## 4. Why MaxPooling?

Reduces computation, shrinks image size, and prevents overfitting.

## 5. Why ReLU?

Prevents vanishing gradients and speeds up training.

## 6. Why Softmax in output?

Softmax converts output logits into probability distribution (sum = 1).

## 7. Why categorical cross-entropy?

Used for multi-class classification with one-hot encoded labels.

## 8. What is normalization?

Scaling pixel values ( $0\text{--}255 \rightarrow 0\text{--}1$ ) for stable training.

## 9. What is backpropagation?

The process of computing gradients and updating weights to reduce loss.

## 10. What is overfitting?

When training accuracy increases but validation accuracy drops.

## ✓ Practical Questions

## 11. Why use ImageDataGenerator?

Automatically loads, resizes, normalizes images.

## 12. Why grayscale?

MNIST images are single-channel.

### 13. Why one-hot encoding?

Required by categorical cross-entropy.

### 14. Why Conv + Pool before Dense?

To extract features and reduce dimensionality.

### 15. What happens if we remove Flatten?

Dense layers cannot process 2D input → error.

### 16. Why train for 5 epochs?

MNIST is simple, and 5 epochs achieve good accuracy.

### 17. Why optimizer='adam'?

Adam adapts learning rate automatically → faster convergence.

### 18. Why batch size = 64?

It balances training speed and stability.

### 19. What does `model.evaluate` do?

Computes loss & accuracy on test data.

### 20. What does `np.argmax` do?

Returns the class with highest predicted probability.

---

Would you like me to create a **1-page viva cheat sheet PDF** for this practical?



14. to summaruze the practical to external explain me in short

Here is the **perfect short viva explanation** you can say to the external. Clear, crisp, and 20–30 seconds — exactly what they expect.

## ✓ Short Summary to Tell the External

"In this practical, I built an image classification model using a CNN on MNIST images. First, I loaded the images from folders using `ImageDataGenerator`, which also resized and normalized them. Then I defined a simple CNN with a `Conv2D` layer to extract features, a `MaxPooling` layer to reduce size, a `Flatten` layer to convert features into a vector, and `Dense` layers for classification using `Softmax`. I trained the model using the `Adam` optimizer for 5 epochs and evaluated it using test data. Finally, I displayed a sample image and compared the actual vs predicted class, and the model achieved good accuracy."

## ✓ Even shorter (15 seconds)

"I loaded MNIST images, normalized them, built a CNN with `Conv2D`, `MaxPooling`, `Flatten`, and `Dense` layers, trained it using `Adam`, evaluated test accuracy, and verified prediction for a sample image."

## ✓ Ultra-short (10 seconds)

"I preprocessed MNIST images, trained a CNN for digit classification, evaluated accuracy, and showed predictions using a simple Conv-Pool-Dense architecture."

If you want, I can also make you a **one-page viva sheet** or **oral script** for flawless delivery.



## 15. ""4ECG

Use Autoencoder to implement anomaly detection. Build the model by using:

- a. Import required libraries
- b. Upload / access the dataset
- c. Encoder converts it into latent representation
- d. Decoder networks convert it back to the original input
- e. Compile the models with Optimizer, Loss, and Evaluation Metrics

```
#Aa. Import required libraries
```

```
import numpy as np
import pandas as pd
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, classification_report
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt
```

```
#B Upload / access the dataset
```

```
# Load the ECG dataset
ecg_dataset = pd.read_csv("ecg.csv")
```

```
# Preprocess the data
```

```
scaler = StandardScaler()
X = scaler.fit_transform(ecg_dataset.values)
y = X # Autoencoder input and output are the same
```

```
# Build and train the Autoencoder model
```

```
input_dim = X_train.shape[1]
```

```
#Cc. Encoder converts it into latent representation
```

```
encoder = models.Sequential([
    layers.Input(shape=(input_dim,)),
    layers.Dense(32, activation='relu'),
    layers.Dense(16, activation='relu'),
    layers.Dense(8, activation='relu')
])
```

```
#DDecoder networks convert it back to the original input
```

```
decoder = models.Sequential([
    layers.Input(shape=(8,)),
    layers.Dense(16, activation='relu'),
    layers.Dense(32, activation='relu'),
    layers.Dense(input_dim, activation='linear') # Use linear activation for reconstruction
])
```

```
#Ee. Compile the models with Optimizer, Loss, and Evaluation Metrics
```

```
autoencoder = models.Sequential([
    encoder,
    decoder
])
autoencoder.compile(optimizer='adam', loss='mean_squared_error')
autoencoder.fit(X_train, X_train, epochs=100, batch_size=32, shuffle=True)
```

```
# Detect anomalies
```

```
y_pred = autoencoder.predict(X_test)
mse = np.mean(np.power(X_test - y_pred, 2), axis=1)
```

```
# Define a threshold for anomaly detection
```

```
threshold = np.percentile(mse, 95) # Adjust the percentile as needed
```

```
# Predict anomalies
anomalies = mse > threshold

# Calculate the number of anomalies
num_anomalies = np.sum(anomalies)
print(f"Number of Anomalies: {num_anomalies}")

# Plot the anomalies
plt.figure(figsize=(12, 6))
plt.plot(mse, marker='o', linestyle='', markersize=3, label='MSE')
plt.axhline(threshold, color='r', linestyle='--', label='Anomaly Threshold')
plt.xlabel('Sample Index')
plt.ylabel('MSE')
plt.title('Anomaly Detection Results')
plt.legend()
plt.show()

plt.figure(figsize=(12, 6))
plt.plot(X_test[0], label='Original ECG')
plt.plot(y_pred[0], label='Reconstructed ECG')
plt.xlabel('Time')
plt.ylabel('Amplitude')
plt.legend()
plt.title('Normal ECG')
plt.show()
```

```
# listing the index of anomalies in X_test
anomalies_index = []
for index, anomaly in enumerate(anomalies):
    if anomaly == True :
        anomalies_index.append(index)

n = 4
anomaly_index = anomalies_index[n]
plt.figure(figsize=(12, 6))
plt.plot(X_test[anomaly_index], label='Original ECG')
plt.plot(y_pred[anomaly_index], label='Reconstructed ECG')
plt.xlabel('Time')
plt.ylabel('Amplitude')
plt.legend()
plt.title('ECG and Anomalies')
plt.show()

# Evaluate the model
y_true = np.zeros(len(X_test))
print("Confusion Matrix:")
print(confusion_matrix(anomalies, anomalies))

print("\nClassification Report:")
print(classification_report(anomalies, anomalies))

import seaborn as sns
```

```
plt.figure(figsize = (6, 4.75))
sns.heatmap(confusion_matrix(anomalies, anomalies), annot = True, annot_kws = {"size": 16},
fmt = 'd')
plt.xticks([0.5, 1.5], rotation = 'horizontal')
plt.yticks([0.5, 1.5], rotation = 'horizontal')
plt.xlabel("Predicted label", fontsize = 14)
plt.ylabel("True label", fontsize = 14)
plt.title("Confusion Matrix", fontsize = 14)
plt.grid(False)
plt.show()
```

for this code explain me this code in three parts for viva first : function by function which does what , important keywords and libraries and third possible questions asked maybe theory based or practical based sppu deep learnig context

And also summarize me the practical to tell external in bit 1-2 minute exokainatin

Below is your full **3-part viva explanation** PLUS a **1-2 minute summary** for external.  
This is crafted exactly in SPPU Deep Learning practical format.

# ✓ PART 1 — FUNCTION-BY-FUNCTION CODE EXPLANATION (VIVA READY)

## ✓ A. Import Required Libraries

python

```
import numpy as np
import pandas as pd
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, classification_report
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt
```

- ✓ **NumPy** → numerical operations
- ✓ **Pandas** → load ECG CSV file
- ✓ **TensorFlow/Keras** → build autoencoder
- ✓ **StandardScaler** → normalize ECG values
- ✓ **confusion\_matrix, classification\_report** → anomaly evaluation
- ✓ **Matplotlib** → plotting ECG + MSE curves

## ✓ B. Upload & Preprocess ECG Dataset

python

```
ecg_dataset = pd.read_csv("ecg.csv")
```

- ✓ Load ECG signals from CSV.

### Normalize / Scale ECG values

python

```
scaler = StandardScaler()
X = scaler.fit_transform(ecg_dataset.values)
y = X
```

- ✓ ECG data is scaled for stable training
- ✓ Autoencoder tries to **reconstruct the same input**, so  
**X = input, y = output**

## ✓ Split Data

(your code uses `X_train`, `X_test` — assumed splitting earlier)

python

```
X_train, X_test = train_test_split(X, test_size=0.2)
input_dim = X_train.shape[1]
```

- ✓ 80% used to train reconstruction
- ✓ 20% used to detect anomalies
- ✓ `input_dim` = number of ECG time steps

## ✓ C. Encoder — Compress Input into Latent Representation

python

```
encoder = models.Sequential([
    layers.Input(shape=(input_dim,)),
    layers.Dense(32, activation='relu'),
    layers.Dense(16, activation='relu'),
    layers.Dense(8, activation='relu')
])
```

- ✓ Autoencoder **compresses ECG signal**
- ✓ Latent dimension = **8**
- ✓ Encoder learns **normal ECG patterns**

## ✓ D. Decoder — Reconstruct Original ECG

python

```
decoder = models.Sequential([
    layers.Input(shape=(8,)),
    layers.Dense(16, activation='relu'),
    layers.Dense(32, activation='relu'),
    layers.Dense(input_dim, activation='linear')
])
```

- ✓ Decoder expands 8-dim latent → reconstruct full ECG
- ✓ Output uses **linear activation** (for continuous values)

## ✓ E. Build Autoencoder

python

```
autoencoder = models.Sequential([encoder, decoder])
autoencoder.compile(optimizer='adam', loss='mean_squared_error')
autoencoder.fit(X_train, X_train, epochs=100, batch_size=32, shuffle=True)
```

- ✓ Autoencoder = Encoder + Decoder
- ✓ **Loss = MSE** → measures reconstruction error
- ✓ **Optimizer = Adam** → fast training
- ✓ Trained using **normal ECG only**

## ✓ F. Detect Anomalies by Reconstruction Error

### Predict reconstructed ECG

```
python
y_pred = autoencoder.predict(X_test)
```

### Calculate Error

```
python
mse = np.mean((X_test - y_pred) ** 2, axis=1)
```

- ✓ Higher MSE → poor reconstruction → likely anomaly

### Set threshold

```
python
threshold = np.percentile(mse, 95)
anomalies = mse > threshold
```

- ✓ Top 5% highest errors marked as anomalies

## ✓ G. Visualizing Results

### Plot MSE + threshold

```
python
plt.plot(mse)
plt.axhline(threshold)
```

- ✓ Shows which samples crossed anomaly threshold

### Plot normal ECG & reconstructed ECG

- ✓ Good match → normal
- ✓ Poor match → anomaly

## Plot confusion matrix

- ✓ Evaluates anomaly detection quality

## ✓ PART 2 — IMPORTANT KEY TERMS & LIBRARIES (VIVA)

### ✓ Important Libraries

Library	Use
Pandas	Load ECG CSV
StandardScaler	Normalize ECG values
Keras layers	Build encoder/decoder
Matplotlib	Plot ECG & error graphs
sklearn.metrics	Confusion matrix

### ✓ Important Keywords

Term	Meaning
<b>Autoencoder</b>	Neural network trained to reconstruct input
<b>Encoder</b>	Reduces data into small latent representation
<b>Decoder</b>	Expands latent vector back to original input
<b>Latent space</b>	The compressed feature vector (8-dim here)
<b>MSE (Mean Squared Error)</b>	Measures reconstruction error
<b>Threshold</b>	If $\text{MSE} > \text{threshold} \rightarrow \text{anomaly}$
<b>Anomaly</b>	ECG that autoencoder cannot reconstruct well

## ✓ PART 3 — LIKELY VIVA QUESTIONS (WITH SHORT ANSWERS)

### ✓ Theory Questions

#### 1. What is an Autoencoder?

A neural network that learns to reconstruct its input using encoder and decoder.

#### 2. Why use autoencoder for anomaly detection?

Because it learns only normal patterns; anomalies reconstruct poorly, giving high error.

#### 3. What is latent representation?

Compressed internal representation of the original signal.

#### 4. Why is MSE used as loss function?

Because it measures difference between original and reconstructed ECG.

## 5. Why linear activation in output layer?

ECG values are continuous, so linear preserves real-valued outputs.

## 6. What is StandardScaler?

It normalizes data to mean=0, variance=1 for stable training.

## Practical Code Questions

### 7. Why X\_train = X\_train in autoencoder training?

The goal is to reconstruct the input, so input = output.

### 8. Why choose threshold = 95 percentile?

Top 5% highest errors often represent anomalies.

### 9. What happens if encoder size is too small?

Model loses information → poor reconstruction.

### 10. What happens if latent dimension is too large?

Model memorizes data → cannot detect anomalies.

### 11. Why shuffle=True?

Prevents model from learning order-based patterns.

## Advanced Questions

### 12. How does autoencoder detect anomalies?

By comparing reconstruction error; high MSE means anomaly.

### 13. Explain backpropagation in autoencoder.

Network adjusts weights to minimize reconstruction error.

### 14. Difference between supervised and unsupervised learning?

Autoencoder is unsupervised because labels are not needed.

### 15. Why Adam optimizer?

It adapts learning rate, converges faster.

## 1-2 MINUTE PRACTICAL SUMMARY FOR EXTERNAL

**"In this practical, I implemented an anomaly detection system for ECG signals using an Autoencoder. First, I loaded the ECG dataset and normalized it using StandardScaler so that the neural network trains properly. Since autoencoders learn to reconstruct the input, I used the same ECG data as both input and output."**

**"Next, I built an autoencoder consisting of an encoder that compresses the ECG signal into an 8-dimensional latent representation, and a decoder that reconstructs the original signal. I trained the model using Adam optimizer and Mean Squared Error loss because we want the reconstruction error to be minimized."**

"After training the model on normal ECG patterns, I passed test data through the autoencoder and calculated reconstruction error (MSE). ECG signals that had very high MSE (above the 95th percentile threshold) were marked as anomalies. I plotted MSE graphs, visualized reconstructed ECGs, displayed anomalous ECGs, and finally evaluated the detection using confusion matrix and classification report."

"This shows how autoencoders can effectively detect abnormal heartbeats by learning only normal patterns and identifying signals that do not match the learned structure."

---

If you want, I can also create:

- A 1-page viva cheat sheet
- A diagram showing encoder-decoder architecture
- A clean formatted practical journal output



16. ""

Use Autoencoder to implement anomaly detection. Build the model by using:

- a. Import required libraries
- b. Upload / access the dataset
- c. Encoder converts it into latent representation
- d. Decoder networks convert it back to the original input
- e. Compile the models with Optimizer, Loss, and Evaluation Metrics

""

#a.Import required libraries

```
import numpy as np
import pandas as pd
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, classification_report
import matplotlib.pyplot as plt
from tensorflow.keras import layers, models
```

#b. Upload / access the dataset

```
dataset = pd.read_csv("creditcard.csv")
```

# Preprocess the data (if needed)

# You may need to normalize the data, handle missing values, etc.

# For simplicity, we'll perform normalization in this example

# Normalize the data to have a mean of 0 and a standard deviation of 1

```
scaler = StandardScaler()
```

```
X = scaler.fit_transform(dataset.drop("Class", axis=1))
```

```
y = dataset["Class"]
```

# Split the dataset into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Build and train the Autoencoder model
input_dim = X_train.shape[1]

#c. Encoder converts it into latent representation
# Encoder
encoder = models.Sequential([
    layers.Input(shape=(input_dim,)),
    layers.Dense(32, activation='relu'),
    layers.Dense(16, activation='relu')
])

#dd. Decoder networks convert it back to the original input
# Decoder
decoder = models.Sequential([
    layers.Input(shape=(16,)),
    layers.Dense(32, activation='relu'),
    layers.Dense(input_dim, activation='linear') # Using 'linear' activation for reconstruction
])

# Autoencoder
autoencoder = models.Sequential([
    encoder,
    decoder
])

#e. Compile the models with Optimizer, Loss, and Evaluation Metrics
autoencoder.compile(optimizer='adam', loss='mean_squared_error')
autoencoder.fit(X_train, X_train, epochs=10, batch_size=32, shuffle=True, validation_data=(X_test, X_test))
```

```
# Detect anomalies and tune the threshold
y_pred = autoencoder.predict(X_test)
mse = np.mean(np.power(X_test - y_pred, 2), axis=1)

# Visualize the reconstruction error distribution
plt.figure(figsize=(10, 6))
plt.hist(mse, bins=50, alpha=0.5, color='b', label='Reconstruction Error')
plt.xlabel("Reconstruction Error")
plt.ylabel("Frequency")
plt.legend()
plt.title("Reconstruction Error Distribution")
plt.show()

# Threshold tuning (iterate and adjust as needed)
thresholds = np.arange(0.1, 1.0, 0.1) # Adjust the step size as needed

for threshold in thresholds:
    anomalies = mse > threshold

# Count the number of anomalies
num_anomalies = np.sum(anomalies)
print(f"Threshold: {threshold:.1f}, Number of anomalies: {num_anomalies}")

# Evaluate the model
print("Confusion Matrix:")
print(confusion_matrix(y_test, anomalies))

print("\nClassification Report:")
print(classification_report(y_test, anomalies))
```

```
import seaborn as sns

plt.figure(figsize = (6, 4.75))
sns.heatmap(confusion_matrix(y_test, anomalies), annot = True, annot_kws = {"size": 16}, fmt =
'd')
plt.xticks([0.5, 1.5], rotation = 'horizontal')
plt.yticks([0.5, 1.5], rotation = 'horizontal')
plt.xlabel("Predicted label", fontsize = 14)
plt.ylabel("True label", fontsize = 14)
plt.title("Confusion Matrix", fontsize = 14)
plt.grid(False)
plt.show()
```

for this code explain me this code in three parts for viva first : function by function which does what , important keywords and libraries and third possible questions asked maybe theory based or practical based sppu deep learnig context And also summarize me the practical to tell external in bit 1-2 minute exokainatin

Below is your complete **3-part viva explanation + a 1-2 minute summary** for external, fully based on **Autoencoder for Credit Card Fraud Detection (Anomaly Detection)**.

This is EXACT, clean, SPPU-style, and ready to speak.

## ✓ PART 1 — FUNCTION-BY-FUNCTION EXPLANATION (Code Walkthrough)

### ✓ A. Import Required Libraries

python

```
import numpy as np
import pandas as pd
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, classification_report
import matplotlib.pyplot as plt
from tensorflow.keras import layers, models
```

What each does:

- **NumPy** → numerical operations
- **Pandas** → load creditcard.csv
- **TensorFlow/Keras** → build autoencoder
- **StandardScaler** → normalize dataset
- **train\_test\_split** → divide into train/test
- **confusion\_matrix, classification\_report** → evaluate anomaly detection
- **Matplotlib** → plot reconstruction error distribution
- **layers, models** → build encoder and decoder

### ✓ B. Load / Access Dataset and Preprocess

## Load dataset

```
python

dataset = pd.read_csv("creditcard.csv")
```

- ✓ Reads the credit-card fraud dataset (highly imbalanced).

## Separate features and labels

```
python

X = scaler.fit_transform(dataset.drop("Class", axis=1))
y = dataset["Class"]
```

- ✓ Normal transactions = 0
- ✓ Fraud transactions = 1
- ✓ StandardScaler normalizes features so neural network trains properly.

## Split into train/test

```
python

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

- ✓ Train on 80% normal data
- ✓ Test on 20% mixed normal + fraudulent transactions

## ✓ C. Build Encoder (Compress High-Dim Data)

```
python

encoder = models.Sequential([
    layers.Input(shape=(input_dim,)),
    layers.Dense(32, activation='relu'),
    layers.Dense(16, activation='relu')
])
```

- ✓ Encoder compresses high-dim credit-card transaction features
- ✓ Learns compressed representation ("latent vector")
- ✓ ReLU adds non-linearity

## ✓ D. Build Decoder (Reconstruct Original Input)

```
python
```

```
decoder = models.Sequential([
    layers.Input(shape=(16,)),
    layers.Dense(32, activation='relu'),
    layers.Dense(input_dim, activation='linear')
])
```

- ✓ Decoder expands latent vector back to original features
- ✓ Output activation = **linear** because input values are continuous

## ✓ Autoencoder = Encoder + Decoder

```
python
```

```
autoencoder = models.Sequential([encoder, decoder])
autoencoder.compile(optimizer='adam', loss='mean_squared_error')
```

- ✓ Loss = **MSE** because we want accurate reconstruction
- ✓ Optimizer = **Adam** → fast, adaptive learning

## ✓ Train Autoencoder

```
python
```

```
autoencoder.fit(X_train, X_train, epochs=10, batch_size=32, shuffle=True, validation_data=(X_test, X_test))
```

- ✓ Autoencoder tries to reconstruct normal transactions
- ✓ Fraudulent transactions → poor reconstruction → higher error (MSE)
- ✓ Input = Output = X (unsupervised learning)

## ✓ Compute Reconstruction Error

```
python
```

```
y_pred = autoencoder.predict(X_test)
mse = np.mean((X_test - y_pred)**2, axis=1)
```

- ✓ MSE measures how well each sample is reconstructed
- ✓ High MSE indicates anomaly/fraud

## ✓ Tune Threshold

python

```
thresholds = np.arange(0.1, 1.0, 0.1)
anomalies = mse > threshold
num_anomalies = np.sum(anomalies)
```

- ✓ Threshold decides if error is high enough to mark anomaly
- ✓ If MSE > threshold → flagged as fraud

## ✓ Evaluate Model

python

```
confusion_matrix(y_test, anomalies)
classification_report(y_test, anomalies)
```

- ✓ Measures accuracy, precision, recall, F1-score
- ✓ Shows model's performance in detecting fraud

## ✓ Visualizations

### Reconstruction error distribution

python

```
plt.hist(mse)
```

### Confusion Matrix

python

```
sns.heatmap(confusion_matrix(y_test, anomalies))
```

- ✓ Helps understand false positives & false negatives

## ✓ PART 2 — IMPORTANT KEYWORDS / LIBRARIES FOR VIVA

## ✓ Important Libraries

- **Pandas** → load CSV
- **StandardScaler** → normalize features
- **Keras layers** → Dense, Input

- **matplotlib/seaborn** → visualize error and confusion matrix
- **sklearn.metrics** → evaluate anomaly detection

## ✓ Important DL Keywords

Term	Meaning
<b>Autoencoder</b>	Unsupervised model that reconstructs input
<b>Encoder</b>	Compresses high-dimensional data
<b>Decoder</b>	Reconstructs original input
<b>Latent space</b>	Compressed representation
<b>MSE</b>	Measures reconstruction error
<b>Reconstruction error</b>	Difference between original and reconstructed
<b>Threshold</b>	Defines anomaly boundary
<b>Anomaly</b>	Sample that reconstruction cannot match → high MSE
<b>Adam optimizer</b>	Adaptive learning rate algorithm

## ✓ PART 3 — POSSIBLE SPPU VIVA QUESTIONS (with answers)

### ✓ Theory Questions

#### 1 What is an autoencoder?

A neural network that learns to reconstruct its input using an encoder and decoder.

#### 2 Why use autoencoder for anomaly detection?

It learns normal patterns; anomalies reconstruct poorly and give high error.

#### 3 What is latent representation?

A compressed low-dimensional encoding of the input features.

#### 4 Why is MSE used as loss?

It measures how different reconstruction is from the original.

#### 5 Why use StandardScaler?

To normalize features for stable and faster training.

## Practical Questions

### 6 Why input and output both are X?

Autoencoder is unsupervised and attempts to reconstruct the same input.

### 7 Why linear activation in final decoder layer?

To output continuous numeric values matching original data.

### 8 What happens if threshold is too low?

Many normal samples are falsely marked as anomalies.

### 9 What if latent dimension is too small?

Reconstruction becomes bad → more false positives.

### 10 Why choose Adam optimizer?

It converges faster and handles sparse, noisy data well.

### 11 Why batch\_size = 32?

Good balance between memory usage and gradient stability.

### 12 Why split into train/test?

Train on normal patterns, test on real data to detect anomalies.

## Advanced Viva Questions

### 13 How does backpropagation work here?

It updates encoder and decoder weights to minimize reconstruction error.

### 14 Why is anomaly detection unsupervised?

Because labels may not be available or fraud cases are very rare.

### 15 What are false positives and false negatives?

FP = normal marked as fraud; FN = fraud missed — critical in banking.

## 1-2 MINUTE PRACTICAL SUMMARY (TELL THIS TO EXTERNAL)

"In this practical, I implemented anomaly detection using an autoencoder on the credit-card fraud dataset. First, I loaded the dataset and normalized all features using StandardScaler so the network trains properly. Since autoencoders learn how to reconstruct normal patterns, I used the transaction data (excluding labels) both as input and output for training."

"The encoder compresses the transaction into a 16-dimensional latent vector, and the decoder reconstructs it back to the original 30 features. I trained the autoencoder using the Adam optimizer and mean squared error loss because reconstruction accuracy is important."

"After training, I passed test data through the model and calculated reconstruction error for each transaction. Normal transactions are reconstructed well, whereas fraudulent ones produce high MSE. Using a threshold, I classified high-error samples as anomalies. Finally, I evaluated performance using confusion matrix and classification report and also plotted MSE distribution and anomaly threshold to visualize results."

"This practical demonstrates how autoencoders can detect outliers or fraud by learning only normal behavior and flagging deviations."

---

If you want, I can also create:

- A 1-page viva-cheat-sheet PDF
- A neat practical journal write-up
- A diagram of the autoencoder architecture



17. 5. Implement the Continuous Bag of Words (CBOW) Model. Stages can be:

- a. Data preparation
- b. Generate training data
- c. Train model
- d. Output

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense,\n    Embedding, Lambda
from tensorflow.keras.preprocessing.text import Tokenizer
import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
import re
```

```
data = """We are about to study the idea of a computational process.\nComputational processes are abstract beings that inhabit computers.\nAs they evolve, processes manipulate other abstract things called data.\nThe evolution of a process is directed by a pattern of rules\ncalled a program. People create programs to direct processes. In effect,\nwe conjure the spirits of the computer with our spells."""
```

```
# for importing data from txt file
# with open("data.txt", "r", encoding="utf-8") as file:
```

```
# data = file.read()
```

```
sentences = data.split(".")
```

```
sentences
```

```
#Clean Data
```

```
clean_sentences = []
```

```
for sentence in sentences:
```

```
    # skip empty string
```

```
    if sentence == "":
```

```
        continue;
```

```
    # remove special characters
```

```
    sentence = re.sub('[^A-Za-z0-9]+', ' ', sentence)
```

```
    # remove 1 letter words
```

```
    sentence = re.sub(r'(?:^ | )\w(?:$ | )', ' ', sentence).strip()
```

```
    # lower all characters
```

```
    sentence = sentence.lower()
```

```
    clean_sentences.append(sentence)
```

```
clean_sentences
```

```
# Define the corpus
corpus = clean_sentences
```

```
# Convert the corpus to a sequence of integers
tokenizer = Tokenizer()
tokenizer.fit_on_texts(corpus)
sequences = tokenizer.texts_to_sequences(corpus)
print("After converting our words in the corpus \
into vector of integers:")
print(sequences)
```

```
# creating dictionary for word to index and index to word
index_to_word_map = {}
word_to_index_map = {}
for index_1, sequence in enumerate(sequences):
    print(sequence)
    words_in_sentence = clean_sentences[index_1].split()
    print(words_in_sentence)
    for index_2, value in enumerate(sequence):
        index_to_word_map[value] = words_in_sentence[index_2]
        word_to_index_map[words_in_sentence[index_2]] = value
```

```
print(index_to_word_map)
print("\n")
print(word_to_index_map)

# Define the parameters
vocab_size = len(tokenizer.word_index) + 1
embedding_size = 10
window_size = 2

# Generate the context-target pairs
contexts = []
targets = []
for sequence in sequences:
    for i in range(window_size, len(sequence) - window_size):
        context = sequence[i - window_size:i] + sequence[i + 1:i + window_size + 1]
        target = sequence[i]
        contexts.append(context)
        targets.append(target)

# sample of training data
```

```
for i in range(5):
    words = []
    target = index_to_word_map.get(targets[i])
    for j in contexts[i]:
        words.append(index_to_word_map.get(j))
    print(words, ">", target)
```

```
# Convert the contexts and targets to numpy arrays
X = np.array(contexts)
Y = np.array(targets)
```

```
# Define the CBOW model
model = Sequential()
model.add(Embedding(input_dim=vocab_size, output_dim=embedding_size, input_length=2 * window_size))
model.add(Lambda(lambda x: tf.reduce_mean(x, axis=1)))
model.add(Dense(256, activation='relu'))
model.add(Dense(512, activation='relu'))
model.add(Dense(units=vocab_size, activation='softmax'))

# Compile the model
model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

# Train the model
model.fit(X, Y, epochs=200, verbose=1)
```

```
# Get the word embeddings
embeddings = model.get_weights()[0]

# Perform PCA to reduce the dimensionality of the embeddings
pca = PCA(n_components=2)
reduced_embeddings = pca.fit_transform(embeddings)

# Visualize the embeddings
plt.figure(figsize=(7, 7))
for i, word in enumerate(tokenizer.word_index.keys()):
    x, y = reduced_embeddings[i]
    plt.scatter(x, y)
    plt.annotate(word, xy=(x, y), xytext=(5, 2),
                 textcoords='offset points',
                 ha='right', va='bottom')
plt.show()

# test model
test_sentences = [
    "we are to study",
    "create programs direct processes",
    "spirits process study program",
    "idea study people create"]
```

]

```
for test_sentence in test_sentences:  
    test_words = test_sentence.split(" ")  
    print("Words: ", test_words)  
    x_test = []  
    for i in test_words:  
        x_test.append(word_to_index_map.get(i))  
    x_test = np.array([x_test])  
    print("Indexes: ", x_test)  
    test_predictions = model.predict(x_test)  
    y_pred = np.argmax(test_predictions[0])  
    print("Predictions: ", test_words, " => ", index_to_word_map.get(y_pred))  
    print("\n")
```

for this code explain me this code in three parts for viva first : function by function which does what , important keywords and libraries and third possible questions asked maybe theory based or practical based sppu deep learnig context And also summarize me the practical to tell external in bit 1-2 minute exokainatin

Below is the **complete CBOW practical viva explanation**, broken into:

- PART 1 → Function-by-function explanation (your code execution flow)
- PART 2 → Important keywords, concepts, libraries
- PART 3 → Viva questions with answers (theory + practical)
- PART 4 → 1-2 minute external exam summary

## PART 1 — FUNCTION-BY-FUNCTION EXPLANATION (Viva-Ready Code Walkthrough)

### A. Import Required Libraries

```
python
```

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Embedding, Lambda
from tensorflow.keras.preprocessing.text import Tokenizer
import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
import re
```

What each does:

- **TensorFlow / Keras** → Build CBOW model
- **Embedding layer** → Word vectors
- **Lambda** → Mean of context word embeddings
- **Tokenizer** → Convert words → integers
- **NumPy** → Arrays
- **PCA** → Reduce embedding dimensions for visualization
- **Matplotlib** → Plot word embeddings
- **re** → Regex cleaning

### B. Data Preparation

Load the text

```
python
```

```
data = """We are about to study..."""
```

Raw text for training-word embeddings.

## Split sentences

```
python
sentences = data.split(".")
```

- ✓ Breaks paragraph into multiple sentences.

## Clean data

```
python
sentence = re.sub('[^A-Za-z0-9]+', ' ', sentence)
sentence = re.sub(r'^(?:(\w+)(\s+))\w(?:(\s+)|$)', '$1$3', sentence)
sentence = sentence.lower()
```

- ✓ Remove special characters
- ✓ Remove single-letter words
- ✓ Convert to lowercase

`clean_sentences` stores cleaned text.

## ✓ C. Tokenization

```
python
tokenizer = Tokenizer()
tokenizer.fit_on_texts(corpus)
sequences = tokenizer.texts_to_sequences(corpus)
```

- ✓ Converts every word into an integer (word → index).
- ✓ `sequences` contains each sentence encoded as numbers.

## Index to Word dictionary

```
python
index_to_word_map[value] = word
word_to_index_map[word] = value
```

- ✓ Helps in reverse mapping for predictions.

## ✓ D. Generate Context-Target Training Data (CBOW Logic)

### Define CBOW parameters

```
python
```

```
embedding_size = 10
window_size = 2
```

- ✓ "Context window" = 2 words before + 2 words after target.

## Generate context-target pairs

```
python
```

```
for sequence in sequences:
    for i in range(window_size, len(sequence) - window_size):
        context = sequence[i-window_size:i] + sequence[i+1:i+1+window_size]
        target = sequence[i]
```

- ✓ Context words = surrounding words

- ✓ Target word = center word

- ✓ Examples:

```
bash
```

```
['we', 'about', 'study', 'idea'] → 'are'
```

## ✓ E. Convert to NumPy Arrays

```
python
```

```
X = np.array(contexts)
Y = np.array(targets)
```

- ✓ X shape = [num\_samples, 4] context words  
 ✓ Y shape = [num\_samples] → target word index

## ✓ F. Build CBOW Model

```
python
```

```
model = Sequential()
model.add(Embedding(input_dim=vocab_size, output_dim=embedding_size, input_length=4))
model.add(Lambda(lambda x: tf.reduce_mean(x, axis=1)))
model.add(Dense(256, activation='relu'))
model.add(Dense(512, activation='relu'))
model.add(Dense(vocab_size, activation='softmax'))
```

## Layer functions:

## ✓ Embedding Layer:

Learns vector of each word (initially random)

## ✓ Lambda Layer:

CBOW → takes **mean** of embeddings of context words

## ✓ Dense Layers:

Deep neural network predicting the target word

## ✓ Softmax Output:

Predicts probability for every word in vocabulary

## ✓ G. Compile & Train

```
python
```

```
model.compile(loss='sparse_categorical_crossentropy', optimizer='adam')  
model.fit(X, Y, epochs=200)
```

- ✓ Loss = sparse cross-entropy (Y is integer)
- ✓ Optimizer = Adam (best for NLP)
- ✓ Trains for 200 epochs

## ✓ H. Extract Embeddings

```
python
```

```
embeddings = model.get_weights()[0]
```

- ✓ First weight matrix in Embedding layer = learned word vectors.

## ✓ I. PCA for Visualization

```
python
```

```
pca = PCA(n_components=2)  
reduced_embeddings = pca.fit_transform(embeddings)
```

- ✓ Reduces word embeddings (10D → 2D) to plot on graph.

## ✓ J. Plot word embeddings

```
python
```

```
plt.scatter(x, y)
plt.annotate(word)
```

- ✓ Each point represents a word in 2D embedding space.

## ✓ K. Test Predictions

```
python
```

```
test_words = "idea study people create"
x_test = [word_to_index_map.get(i)]
test_predictions = model.predict(x_test)
```

- ✓ Model predicts center word from context.  
 ✓ Demonstrates CBOW functionality.

## ✓ PART 2 — IMPORTANT KEYWORDS & LIBRARIES (VIVA)

### ✓ Keywords

Term	Meaning
<b>CBOW</b>	Predict center word from context
<b>Word Embedding</b>	Numerical vector representation of words
<b>Context Window</b>	Surrounding words used for prediction
<b>Tokenizer</b>	Converts text → integers
<b>Softmax</b>	Gives probability for each word
<b>Sparse Categorical Crossentropy</b>	Loss for integer-coded labels
<b>PCA</b>	Dimensionality reduction
<b>Embedding Layer</b>	Learns vector representation of each word

### ✓ Libraries

- **TensorFlow/Keras** → Model building
- **Tokenizer** → Preprocess text
- **PCA** → Reduce embedding dimensions
- **Matplotlib** → Plot word vectors
- **Regex** → Clean text

- NumPy → Construct training samples

## ✓ PART 3 — Likely SPPU Viva Questions (With Answers)

### ✓ Theory Questions

#### ✓ 1. What is CBOW?

CBOW predicts the **target word** using its **surrounding context words**.

#### ✓ 2. Difference between CBOW and Skip-gram?

CBOW: Context → Target

Skip-gram: Target → Predict context words

#### ✓ 3. What are word embeddings?

Dense vector representations that capture semantic meaning of words.

#### ✓ 4. What is softmax used for?

Softmax outputs probability distribution across all vocabulary words.

#### ✓ 5. Why use sparse categorical cross-entropy?

Because target labels are integer IDs, not one-hot vectors.

#### ✓ 6. Why use PCA?

To reduce multi-dimensional embeddings to 2D for visualization.

### ✓ Practical Questions

#### ✓ 7. Why use Tokenizer?

To convert words into numeric indices that the model can process.

#### ✓ 8. Why clean the text?

To remove special characters, lower-case, remove garbage words.

#### ✓ 9. Why use Embedding layer?

It learns vector representation of words during training.

## ✓ 10. Why take mean of context embeddings?

Core concept of CBOW → average context represents meaning.

## ✓ 11. What is window size?

Number of words taken before and after the target.

## ✓ 12. What does context-target pair mean?

Context = surrounding words

Target = word in the center

## ✓ PART 4 — 1-2 Minute External Exam Summary (Say This)

**"In this practical, I implemented the Continuous Bag of Words (CBOW) model for word embeddings. I began by loading a paragraph of text and cleaning it using regex by removing special characters, converting to lowercase, and removing single-letter words. Next, I tokenized the words into integer IDs and created context-target pairs using a window size of 2, where the surrounding 4 words become context and the middle word becomes the target."**

**"I built the CBOW neural network using an Embedding layer to learn word vectors, a Lambda layer to average the embeddings of context words, two Dense layers for deeper learning, and a Softmax output layer to predict the target word. The model was compiled using sparse categorical cross-entropy and trained for 200 epochs."**

**"After training, I extracted the learned word embeddings and used PCA to reduce them to 2D for visualization. Finally, I tested the model with custom sentences to predict the missing center word. This practical demonstrates how CBOW learns meaningful word embeddings from raw text."**

If you want:

- ✓ A one-page viva cheat sheet
- ✓ A colorful CBOW architecture diagram
- ✓ A handwritten-style journal summary

Just tell me!



## 18. 6. Object detection using Transfer Learning of CNN architectures

- a. Load in a pre-trained CNN model trained on a large dataset
- b. Freeze parameters (weights) in model's lower convolutional layers
- c. Add custom classifier with several layers of trainable parameters to model
- d. Train classifier layers on training data available for task
- e. Fine-tune hyper parameters and unfreeze more layers as needed

dataset: CIFAR

```
import tensorflow as tf
from tensorflow import keras
```

```
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import numpy as np
```

```
train_dir = "cifar-10-img/train"
test_dir = "cifar-10-img/test"
```

```
train_datagen = ImageDataGenerator(
    rescale=1.0 / 255,
)
```

```
test_datagen = ImageDataGenerator(
    rescale=1.0 / 255,
)
```

```
# here batch_size is the number of images in each batch
```

```
train_batch_size = 5000
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(32, 32),
    batch_size=train_batch_size,
    class_mode='categorical'
)
test_batch_size = 1000
test_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=(32, 32),
    batch_size=test_batch_size,
    class_mode='categorical'
)

x_train, y_train = train_generator[0]
x_test, y_test = test_generator[0]

print(len(x_train))
print(len(x_test))

# Load VGG16 without top layers
weights_path = "vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5"
base_model = VGG16(weights=weights_path, include_top=False, input_shape=(32, 32, 3))

for layer in base_model.layers:
```

```
layer.trainable = False
```

```
x = Flatten()(base_model.output)
x = Dense(256, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

# Create the model
model = Model(inputs=base_model.input, outputs=predictions)
# Compile the model
model.compile(optimizer="adam", loss='categorical_crossentropy', metrics=['accuracy'])
```

```
# Train the model
model.fit(x_train, y_train, batch_size=64, epochs=10, validation_data=(x_test, y_test))
```

```
base_model = VGG16(weights=weights_path, include_top=False, input_shape=(32, 32, 3))
# freeze all layers first
for layer in base_model.layers:
    layer.trainable = False
# unfreeze last 4 layers of base model
for layer in base_model.layers[len(base_model.layers) - 4:]:
    layer.trainable = True
```

```
# fine-tuning hyper parameters
x = Flatten()(base_model.output)
x = Dense(256, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

# Create the model
model = Model(inputs=base_model.input, outputs=predictions)
# Compile the model
model.compile(optimizer=Adam(learning_rate=0.001), loss='categorical_crossentropy', metrics=['accuracy'])
# training fine tuned model
model.fit(x_train, y_train, batch_size=64, epochs=10, validation_data=(x_test, y_test))

import matplotlib.pyplot as plt
predicted_value = model.predict(x_test)

labels = list(test_generator.class_indices.keys())

n = 756
plt.imshow(x_test[n])
print("Predicted: ", labels[np.argmax(predicted_value[n])])
print("Actual: ", labels[np.argmax(y_test[n])])
```

for this code explain me this code in three parts for viva first : function by function in answer also tell something about function which does what , important keywords and libraries and third possible questions asked maybe theory based or practical based sppu deep learnig context And also summarize me the practical to tell external in bit 1-2 minute exokainatin

Below is the **complete, SPPU-style, perfect viva explanation** for your practical:

- PART 1 — Function-by-function code explanation**
- PART 2 — Important keywords, concepts, libraries**
- PART 3 — Viva questions (theory + practical + advanced)**
- PART 4 — 1-2 minute external exam summary**

This will fully prepare you for your viva.

## **PART 1 — FUNCTION-BY-FUNCTION CODE EXPLANATION (Viva Ready)**

### **A. Import Libraries**

```
python
```

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import numpy as np
```

**What each does:**

- **VGG16** → Pretrained CNN model (on ImageNet)
- **Model** → Functional API for modifying pretrained models
- **Dense, Flatten** → Custom classifier layers
- **ImageDataGenerator** → Load CIFAR-10 images from folders
- **Adam** → Optimizer for training
- **NumPy** → Numerical operations

### **B. Load and Preprocess CIFAR-10 Dataset**

## Create data generators

python

```
train_datagen = ImageDataGenerator(rescale=1.0/255)
test_datagen = ImageDataGenerator(rescale=1.0/255)
```

- ✓ Normalizes pixel values
- ✓ Essential for CNN training

## Load images from directory

python

```
train_generator = train_datagen.flow_from_directory(...)
test_generator = test_datagen.flow_from_directory(...)
```

- ✓ Reads images
- ✓ Resizes to 32x32
- ✓ One-hot encodes labels (`class_mode='categorical'`)
- ✓ Loads entire dataset in one batch (5000 train, 1000 test)

## Extract actual data

python

```
x_train, y_train = train_generator[0]
x_test, y_test = test_generator[0]
```

- ✓ Because batch size = dataset size
- ✓ Index 0 gives the entire dataset

## ✓ C. Load Pretrained VGG16 (Transfer Learning)

python

```
base_model = VGG16(weights=weights_path, include_top=False, input_shape=(32,32,3))
```

- ✓ Loads VGG16 WITHOUT classifier head
- ✓ Uses pretrained convolution layers
- ✓ Pretrained on ImageNet → millions of images
- ✓ Keeps its knowledge about edges, textures, shapes

## ✓ D. Freeze Pretrained Layers

```
python
```

```
for layer in base_model.layers:  
    layer.trainable = False
```

- ✓ Freezes lower convolution layers
- ✓ Prevents updating pretrained weights
- ✓ We only train our custom classifier on top

## ✓ E. Add Custom Classifier

```
python
```

```
x = Flatten()(base_model.output)  
x = Dense(256, activation='relu')(x)  
predictions = Dense(10, activation='softmax')(x)
```

- ✓ Flatten → convert feature maps into 1D vector
- ✓ Dense(256) → learn patterns specific to CIFAR-10
- ✓ Dense(10, softmax) → 10 output classes (CIFAR-10 dataset)

## Create & Compile final model

```
python
```

```
model = Model(inputs=base_model.input, outputs=predictions)  
model.compile(optimizer="adam", loss='categorical_crossentropy', metrics=['accuracy'])
```

- ✓ Adam → adaptive, fast
- ✓ Loss → multi-class classification
- ✓ Tracks accuracy

## ✓ F. Train Custom Classifier

```
python
```

```
model.fit(x_train, y_train, batch_size=64, epochs=10, validation_data=(x_test, y_test))
```

- ✓ Only classifier layers train
- ✓ Pretrained layers remain unchanged

## ✓ G. Fine-Tuning (Advanced Transfer Learning)

### Reload base model

```
python
```

```
base_model = VGG16(...include_top=False...)
```

## Freeze all layers first

```
python
```

```
for layer in base_model.layers:  
    layer.trainable = False
```

## Unfreeze last 4 layers only

```
python
```

```
for layer in base_model.layers[-4:]:  
    layer.trainable = True
```

- ✓ Allows deeper layers to adjust to CIFAR-10 dataset
- ✓ Fine-tuning improves accuracy

## Rebuild classifier and train again

```
python
```

```
model.compile(optimizer=Adam(learning_rate=0.001), ...)  
model.fit(...)
```

- ✓ Smaller learning rate for fine-tuning

## ✓ H. Prediction

```
python
```

```
predicted_value = model.predict(x_test)  
labels = list(test_generator.class_indices.keys())  
plt.imshow(x_test[n])  
print("Predicted: ", labels[argmax])
```

- ✓ Shows actual vs predicted class
- ✓ Visual verification of model performance

## ✓ PART 2 — IMPORTANT KEYWORDS & LIBRARIES

## ✓ Important Concepts

Term	Meaning
<b>Transfer Learning</b>	Using knowledge from pre-trained CNN models
<b>Feature Extraction</b>	Using pretrained layers as filters
<b>Freezing layers</b>	Prevents updating pretrained weights
<b>Fine-tuning</b>	Slightly retraining deeper layers
<b>Flatten</b>	Convert 3D → 1D for dense layers
<b>Softmax</b>	Converts outputs to probabilities
<b>Categorical crossentropy</b>	Loss for multi-class classification

## ✓ Important Libraries

- **keras.applications.VGG16** → pretrained model
- **ImageDataGenerator** → image loading & preprocessing
- **Model (Functional API)** → modify pretrained models
- **Adam optimizer** → training

## ✓ PART 3 — POSSIBLE VIVA QUESTIONS (SPPU) + ANSWERS

### ✓ Theory Questions

#### 1. What is transfer learning?

Using a pretrained CNN model trained on a large dataset and adapting it to a new task.

#### 2. Why freeze layers?

To preserve pretrained knowledge and avoid retraining huge networks.

#### 3. Why unfreeze last layers during fine-tuning?

Because deeper layers learn high-level features more relevant to the new dataset.

#### 4. What is VGG16?

A pretrained CNN architecture with 16 layers trained on ImageNet.

#### 5. Why use softmax in output?

Softmax gives probability distribution across 10 classes.

## ✓ Practical Questions

### 6. Why target\_size=32x32?

CIFAR-10 images are originally 32×32.

### 7. Why categorical class mode?

Because CIFAR-10 has 10 classes (multi-class classification).

### 8. Why include\_top=False?

To remove ImageNet classifier and add custom classifier.

### 9. Why use Adam optimizer?

Adaptive learning rate → trains faster.

### 10. Why convert the feature maps to 1D?

Dense classifier only accepts 1D inputs → done using Flatten.

### 11. Why load full dataset with a large batch?

Because CIFAR-10 is small and fits in memory.

## ✓ Advanced Questions

### 12. What happens if you do not freeze layers?

The model may overfit or destroy pretrained knowledge.

### 13. Difference between feature extraction and fine-tuning?

Feature extraction = freeze all layers

Fine-tuning = unfreeze some deeper layers

### 14. Why learning rate = 0.001 in fine-tuning?

Small LR prevents damaging pretrained weights.

### 15. Why VGG16 is good for transfer learning?

It has strong general-purpose filters trained on millions of images.

## ✓ PART 4 — 1-2 MINUTE SUMMARY FOR EXTERNAL

"In this practical, I implemented object classification using transfer learning. I used the CIFAR-10 dataset and a pretrained VGG16 model, which was originally trained on ImageNet. First, I loaded and normalized the CIFAR images using ImageDataGenerator. Then I loaded VGG16 with include\_top=False to remove its classifier, keeping only the convolutional feature extractor. I froze its lower layers to preserve pretrained weights and added my own custom classifier consisting of Flatten, Dense(256), and Dense(10) with Softmax."

"I trained only the custom classifier first, treating VGG16 as a fixed feature extractor. After initial training, I fine-tuned the model by unfreezing the last four layers of VGG16 with a small learning rate so the model could adapt better to CIFAR-10 features. Finally, I evaluated the model on test images and displayed predictions along with actual labels. This demonstrates how transfer learning reduces training time and improves accuracy by reusing knowledge from pretrained CNNs."

---

If you want, I can also provide:

- ✓ A one-page viva cheat sheet
- ✓ Diagram of transfer learning architecture
- ✓ Journal-ready practical writeup



## 19. 6. Object detection using Transfer Learning of CNN architectures

- a. Load in a pre-trained CNN model trained on a large dataset
- b. Freeze parameters (weights) in model's lower convolutional layers
- c. Add custom classifier with several layers of trainable parameters to model
- d. Train classifier layers on training data available for task
- e. Fine-tune hyper parameters and unfreeze more layers as needed

dataset: Caltech

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import numpy as np
```

```
dataset_dir = "caltech-101-img/"
dataset_datagen = ImageDataGenerator(
    rescale=1.0 / 255,
)

# here batch_size is the number of images in each batch
batch_size = 2000
dataset_generator = dataset_datagen.flow_from_directory(
    dataset_dir,
    target_size=(64, 64),
    batch_size=batch_size,
    class_mode='categorical'
)
```

```
x_train, y_train = dataset_generator[0]
x_test, y_test = dataset_generator[1]

print(len(x_train))
print(len(x_test))

# Load VGG16 without top layers
weights_path = "vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5"
base_model = VGG16(weights=weights_path, include_top=False, input_shape=(64, 64, 3))

for layer in base_model.layers:
    layer.trainable = False

x = Flatten()(base_model.output)
x = Dense(64, activation='relu')(x)
predictions = Dense(102, activation='softmax')(x)

# Create the model
model = Model(inputs=base_model.input, outputs=predictions)
# Compile the model
model.compile(optimizer="adam", loss='categorical_crossentropy', metrics=['accuracy'])
```

```
# Train the model  
model.fit(x_train, y_train, batch_size=64, epochs=10, validation_data=(x_test, y_test))
```

```
import matplotlib.pyplot as plt  
predicted_value = model.predict(x_test)
```

```
labels = list(dataset_generator.class_indices.keys())
```

```
n = 1000  
plt.imshow(x_test[n])  
print("Predicted: ", labels[np.argmax(predicted_value[n])])  
print("Actual: ", labels[np.argmax(y_test[n])])
```

for this code explain me this code in three parts for viva first : function by function in answer also tell something about function which does what , important keywords and libraries and third possible questions asked maybe theory based or practical based sppu deep learnig context And also summarize me the practical to tell external in bit 1-2 minute exokainatin

Below is the **complete SPPU-ready viva explanation** for your **Transfer Learning – VGG16 on Caltech-101 dataset**, broken into:

- PART 1 — Function-by-function explanation (with meaning of each function)**
- PART 2 — Important keywords, layers, libraries**
- PART 3 — Likely SPPU viva questions with answers**
- PART 4 — 1-2 minute external exam summary**

Use this directly in your practical/viva.

## **PART 1 — FUNCTION-BY-FUNCTION CODE EXPLANATION (Viva-Ready)**

### **A. Import Required Libraries**

```
python

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import numpy as np
```

#### **What each does:**

- **VGG16** → Pretrained CNN backbone (from ImageNet)
- **Model (Functional API)** → Used to build a new network around VGG16
- **Flatten** → Converts feature maps → vector
- **Dense** → Custom classifier layers
- **ImageDataGenerator** → Loads and preprocesses Caltech images
- **Adam optimizer** → Used to train the custom classifier
- **NumPy** → Numerical processing

### **B. Load and Preprocess Caltech-101 Dataset**

#### **ImageDataGenerator**

```
python

dataset_datagen = ImageDataGenerator(rescale=1.0 / 255)
```

- ✓ Normalizes image pixel values (0–255 → 0–1).
- ✓ Improves training stability.

## Load Dataset

```
python

dataset_generator = dataset_datagen.flow_from_directory(
    dataset_dir,
    target_size=(64, 64),
    batch_size=2000,
    class_mode='categorical'
)
```

- ✓ Reads images directly from folders
- ✓ Resizes all images to 64×64×3
- ✓ Converts labels to one-hot encoding
- ✓ Loads all images in large batches (Caltech is small)

## Split into Train/Test

```
python

x_train, y_train = dataset_generator[0]
x_test, y_test = dataset_generator[1]
```

- ✓ Since batch size = 2000, index 0 and 1 give two different random batches
- ✓ First batch used as train data
- ✓ Second batch used as test data

## ✓ C. Load Pre-Trained VGG16 Model

```
python

base_model = VGG16(weights=weights_path, include_top=False, input_shape=(64, 64, 3))
```

- ✓ Loads VGG16 **without the top fully connected layers**
- ✓ Includes only convolutional feature extractor
- ✓ Pretrained on ImageNet → 1.2 million images

## ✓ D. Freeze Base Model Layers

```
python

for layer in base_model.layers:
```

```
layer.trainable = False
```

- ✓ Freezes all pretrained convolution layers
- ✓ Prevents changing the weights
- ✓ Uses VGG16 as a **feature extractor**

## ✓ E. Add Custom Classifier

python

```
x = Flatten()(base_model.output)
x = Dense(64, activation='relu')(x)
predictions = Dense(102, activation='softmax')(x)
```

- ✓ Flatten → converts spatial features to 1D
- ✓ Dense(64) → learns Caltech-specific patterns
- ✓ Dense(102, softmax) → because Caltech-101 has **102 class folders**

## ✓ F. Define and Compile Final Model

python

```
model = Model(inputs=base_model.input, outputs=predictions)
model.compile(optimizer="adam", loss='categorical_crossentropy', metrics=['accuracy'])
```

- ✓ Adam → fast adaptive optimizer
- ✓ Loss → multi-class classification
- ✓ Metric → accuracy

## ✓ G. Train Only Custom Classifier Layers

python

```
model.fit(x_train, y_train, batch_size=64, epochs=10, validation_data=(x_test, y_test))
```

- ✓ Only Dense layers train
- ✓ VGG16 layers remain frozen
- ✓ This is **transfer learning stage-1**

## ✓ H. Make Predictions

python

```
predicted_value = model.predict(x_test)
labels = list(dataset_generator.class_indices.keys())
```

- ✓ Predict class probabilities
- ✓ Decode from index → class name

## Show one image and prediction

python

```
plt.imshow(x_test[n])
print("Predicted:", labels[np.argmax(predicted_value[n])])
print("Actual:", labels[np.argmax(y_test[n])])
```

- ✓ Visual check of model accuracy

## ✓ PART 2 — IMPORTANT KEYWORDS & LIBRARIES

### ✓ Important Concepts

Term	Meaning
<b>Transfer Learning</b>	Using pretrained CNN for a new task
<b>Feature Extraction</b>	Using frozen convolution layers
<b>Fine-tuning</b>	Unfreezing some layers to retrain them
<b>Weights freezing</b>	Locking pretrained weights
<b>Pretrained Model</b>	A CNN trained on a large dataset
<b>Softmax</b>	Probability distribution for 102 classes
<b>Categorical Crossentropy</b>	Loss for multi-class classification
<b>Flatten</b>	Converts feature map → 1D vector

### ✓ Important Libraries

Library	Use
<b>VGG16</b>	Pretrained CNN backbone
<b>ImageDataGenerator</b>	Load and preprocess images
<b>Model (Functional API)</b>	Connect frozen and custom layers
<b>Adam optimizer</b>	Train classifier layers

## ✓ PART 3 — SPPU VIVA QUESTIONS WITH ANSWERS

### ✓ Theory Questions

#### 1. What is transfer learning?

Using a pretrained CNN model trained on a large dataset like ImageNet and adapting it to a smaller dataset.

#### 2. Why freeze layers?

To preserve pretrained features like edges, textures, patterns.

#### 3. What is VGG16?

A deep CNN architecture with 16 layers trained on ImageNet.

#### 4. What is feature extraction?

Using pretrained convolution layers to extract useful features.

#### 5. Why use softmax?

Softmax outputs probabilities for 102 classes (Caltech-101).

#### 6. Why use categorical\_crossentropy?

It is the standard loss function for multi-class classification.

### ✓ Practical Questions

#### 7. Why include\_top=False?

To remove default classifier and attach your own.

#### 8. Why train only Dense layers first?

To train quickly without damaging pretrained weights.

#### 9. Why resize images to 64x64?

VGG16 requires consistent input size; Caltech images vary in size.

#### 10. Why use ImageDataGenerator?

Loads images from folders, normalizes, batches automatically.

## 11. What happens if you unfreeze too many layers?

Model may overfit or destroy pretrained knowledge.

## 12. Why batch\_size=2000?

Caltech dataset is small; loading all images at once is efficient.

## Advanced Questions

### 13. What is fine-tuning?

Retraining selected deeper layers of pretrained model to adapt it better to the new dataset.

### 14. Why use Adam?

Adaptive learning rate, efficient on deep networks.

### 15. Difference between CNN and pretrained CNN?

A pretrained CNN already has learned features from a huge dataset.

## PART 4 — 1-2 MINUTE PRACTICAL SUMMARY (Tell This to External)

**“In this practical, I implemented object classification using transfer learning on the Caltech-101 dataset. I used VGG16 as a pretrained CNN model trained on ImageNet. First, I loaded and normalized images using ImageDataGenerator, and resized them to 64×64. I loaded the VGG16 model without the top classifier layers, keeping only the convolutional backbone. These layers were frozen so that their pretrained weights remain unchanged.”**

**“Next, I added my own custom classifier: a Flatten layer, a Dense layer with 64 neurons, and a final Dense layer with 102 softmax outputs for the Caltech-101 classes. I trained only the new classifier layers while keeping the base VGG16 frozen. After training, I evaluated the model on test images and visualized predictions. This practical demonstrates how to reuse a powerful CNN model for a new image-classification task with much less data and training time.”**

If you want, I can also prepare:

- A one-page viva answer sheet
- A neat practical write-up for journal
- A diagram of transfer learning architecture